

FINAL_PROJECT

August 12, 2020

0.1 # Final Project

0.2 ## State of the art CNN architecture over imagenet dataset

- Duke Mervyn Martin (dukemerv@buffalo.edu)
- Preeti Kumari (preetiku@buffalo.edu)
- Kizito Nwaka (kizitonw@buffalo.edu)

```
[1]: import tensorflow as tf
from tensorflow import keras
from keras.models import Sequential
from keras.layers import Dense, MaxPooling2D, Flatten, Conv2D, LeakyReLU,
↳Activation, BatchNormalization, Dropout
from keras import optimizers, regularizers
from keras.utils import plot_model
from keras.optimizers import SGD
from keras import backend
import numpy as np
import time as t
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
import os
import sys
os.environ["CUDA_VISIBLE_DEVICES"]="0"
#tf.debugging.set_log_device_placement(True)
%matplotlib inline
import seaborn as sns
import warnings
warnings.filterwarnings(action='once')
import cv2
import matplotlib.image as mpimg
from keras.utils import plot_model

[2]: print(f"Tensor Flow Version: {tf.__version__}")
print(f"Keras Version: {keras.__version__}")
print("GPU is", "available" if tf.test.is_gpu_available() else "NOT AVAILABLE")
print("Num GPUs Available: ", len(tf.config.experimental.
↳list_physical_devices('GPU')))
```

Tensor Flow Version: 2.3.0

Keras Version: 2.4.0

WARNING:tensorflow:From <ipython-input-2-227a307d88ee>:3: is_gpu_available (from tensorflow.python.framework.test_util) is deprecated and will be removed in a future version.

Instructions for updating:

Use ``tf.config.list_physical_devices('GPU')`` instead.

GPU is available

Num GPUs Available: 1

0.2.1 Choosing classes Randomly and Downloading 1500 Images per class

From the library [ImageNet-Datasets-Downloader](#) we use the `downloader.py` with the command given below to download all the Images respective to its classes in each folder inside the directory `imagenet/`

```
python ./downloader.py -data_root imagenet -number_of_classes 8 -images_per_class 1500
```

List of Classes

```
[4]: from os import listdir
      classes = listdir("imagenet")
      print('List of Classes ->', classes)
      print("Number of classes =", len(classes))
```

List of Classes -> ['koala', 'ladybug', 'lichen', 'meerkat', 'Rhodesian ridgeback', 'tamandua', 'vizsla', 'Yorkshire terrier']

Number of classes = 8

0.2.2 Preparing the Datasets

As all the images are of different sizes, we will be generating a square image of size **n x n**

This can be achieved by two steps - Scale down the image to n pixels by its shortest side - Crop the middle part of the image to get n x n pixels square (This method of data preparation is used by AlexNet model, VGG-16 & VGG-19)

```
[2]: n = 128
```

0.3 Image process demo


```
[5]: img = cv2.imread('imagenet/koala/218758752_650f6b9b5a.jpg')
      print('Image Before Processing')
      print('-----')
      plt.imshow(img)
      plt.show()
      print(img.shape)
      print('Image After Resizing')
      print('-----')
      if(img.shape[0]>img.shape[1]):
```

```

width = n
height = int(img.shape[0] * (n/img.shape[1]))
img = cv2.resize(img,(width,height))
plt.imshow(img)
plt.show()
print(img.shape)
else:
width = int(img.shape[1] * (n/img.shape[0]))
height = n
img = cv2.resize(img,(width,height))
plt.imshow(img)
plt.show()
print(img.shape)
print('Image After Cropping')
print('-----')
if(img.shape[0]==n):
x = int((img.shape[1]-n)/2)
img = img[0:,x:x+n]
plt.imshow(img)
print(img.shape)
else:
x = int((img.shape[0]-n)/2)
img = img[x:x+n,0:]
plt.imshow(img)
plt.show()
print(img.shape)

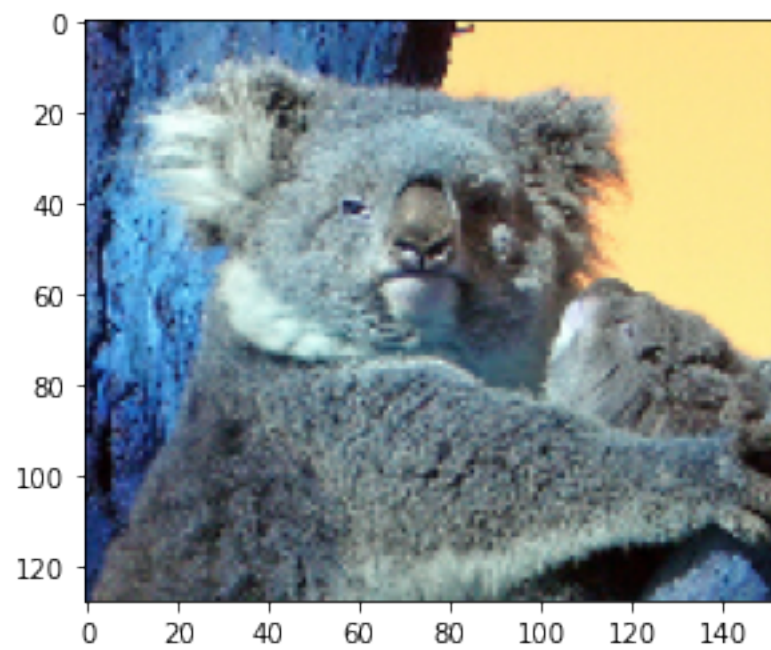
```

Image Before Processing



(416, 500, 3)

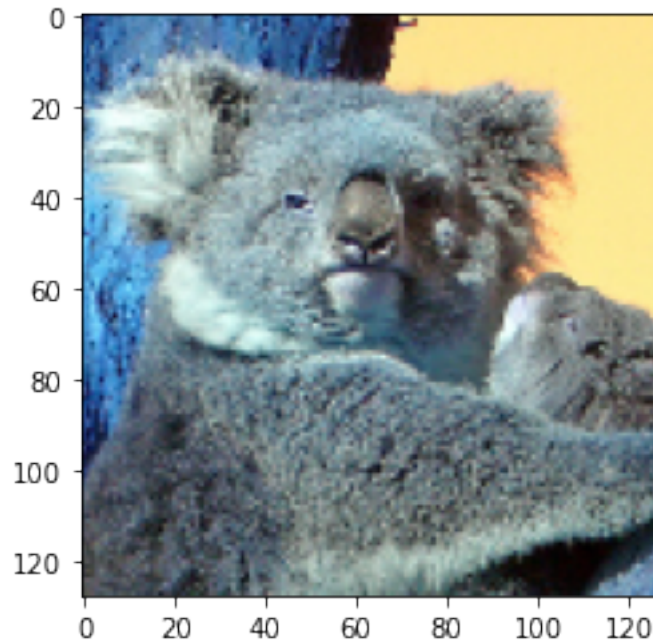
Image After Resizing



(128, 153, 3)

Image After Cropping

(128, 128, 3)



Preprocessing Images

```
[4]: def img_process(image):  
    img = image  
    if(img.shape[0]>img.shape[1]):  
        width = n  
        height = int(img.shape[0] * (n/img.shape[1]))  
        img = cv2.resize(img,(width,height))  
    else:  
        width = int(img.shape[1] * (n/img.shape[0]))  
        height = n  
        img = cv2.resize(img,(width,height))  
  
    if(img.shape[0]==n):  
        x = int((img.shape[1]-n)/2)  
        img = img[0:,x:x+n]  
    else:  
        x = int((img.shape[0]-n)/2)  
        img = img[x:x+n,0:]  
    img = cv2.resize(img,(n,n))  
    return img
```

Generate Target Labels

```
[7]: labels= [] # Target Classes
    for name in classes:
        for jpg in listdir("imagenet/"+name):
            labels.append([name])
```

Generate Colour Image Datasets after applying the preprocessing function

```
[8]: labels=np.array(labels)
```

```
[9]: labels.shape
```

```
[9]: (12000, 1)
```

```
[10]: images = np.zeros((labels.shape[0],n,n,3)) #Empty Dataset
    i=0
    for name in classes:
        for jpg in listdir("imagenet/"+name):
            img = cv2.imread('imagenet/'+name+'/'+jpg)
            img = img_process(img)
            images[i] = img
            i+=1
    print('Image data shape ->',images.shape)
```

Image data shape -> (12000, 128, 128, 3)

```
[11]: # Save the Data for future use
    np.save('image_data_128_color',images)
    np.save('image_label_128_color',labels)
```


```
[4]: image_data= np.load('image_data_128_color.npy')
    labels= np.load('image_label_128_color.npy')
```

```
[5]: print('Check Max Min before Normalization ->',image_data.min(),image_data.max())
    image_data = image_data/255
    print('Check Max Min after Normalization ->',image_data.min(),image_data.max())
```

Check Max Min before Normalization -> 0.0 255.0

Check Max Min after Normalization -> 0.0 1.0

```
[6]: plt.figure(figsize=(15,15))
    for i in range(5):
        plt.subplot(5,5,i+1)
        plt.imshow(image_data[i])
        plt.xlabel(labels[i])
    plt.show()
```



0.3.1 Split Training and Testing Data

```
[8]: from sklearn.utils import shuffle

image_data, labels = shuffle(image_data, labels)
np.save('shuffled_images_128', image_data)
np.save('shuffled_label_128', labels)
```

- We use a common shuffled data for analysing the performance of all the models, thus we reduce the effect of shuffling attributing to the accuracy of the model.

```
[4]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder
```

```
c:\users\dm97o\anaconda3\envs\tf_gpu\lib\importlib\_bootstrap.py:219:
RuntimeWarning: numpy.ufunc size changed, may indicate binary incompatibility.
Expected 192 from C header, got 216 from PyObject
    return f(*args, **kwds)
```


```
[5]: images= np.load('shuffled_images_128.npy')
labels= np.load('shuffled_label_128.npy')
```

- We load the shuffled data and encode the labels using the onehot encoder.

```
[6]: label_encoder = OneHotEncoder()
label_encoder.fit(labels)
label_encoded = label_encoder.transform(labels).toarray()
train_images, test_images = images[0:10000], images[10000:]
train_labels, test_labels = label_encoded[0:10000], label_encoded[10000:]
```

0.3.2 Some statistics about the data

```
[7]: plt.figure(figsize=(15,15))
      for i in range(5):
          plt.subplot(5,5,i+1)
          plt.imshow(train_images[i])
          plt.xlabel(labels[i])
      plt.show()
```



```
[8]: print(labels[:5])
      print(label_encoded[:5])
```

```
['Rhodesian ridgeback']
['ladybug']
['ladybug']
['tamandua']
['vizsla']
[[1. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 1. 0. 0. 0. 0.]
 [0. 0. 0. 1. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 1. 0.]
 [0. 0. 0. 0. 0. 0. 0. 1.]]
```

```
[9]: print('Total Number of Entries ->',len(labels))
      print('X Train Shape ->',train_images.shape)
      print('Y Train Shape ->',train_labels.shape)
      print('X Test Shape ->',test_images.shape)
      print('Y Test Shape ->',test_labels.shape)
```

```
Total Number of Entries -> 12000
X Train Shape -> (10000, 128, 128, 3)
Y Train Shape -> (10000, 8)
X Test Shape -> (2000, 128, 128, 3)
Y Test Shape -> (2000, 8)
```

```
[8]: def graphs(history,score):

      plt.figure(figsize=(15,7))
      plt.plot(history.history['accuracy'])
      plt.plot(history.history['val_accuracy'])
```



```
plt.title('Model accuracy',size=20)
plt.ylabel('Accuracy',size=15)
plt.xlabel('Epochs', size=15)
plt.legend(['Training', 'Testing'], loc='best')
plt.show()

plt.figure(figsize=(15,7))
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss',size=20)
plt.ylabel('Loss',size=15)
plt.xlabel('Epochs',size=15)
plt.legend(['Training', 'Testing'], loc='best')
plt.show()

print("Accuracy of the model: {:.2f}%".format(score[1]*100))
```

0.3.3 Building the network

```
[9]: opt = SGD(lr=0.01,momentum=0.9,decay=0.01)
```

```
[10]: from keras.callbacks import ReduceLROnPlateau, EarlyStopping
early = EarlyStopping(monitor='val_loss', min_delta=0, patience=10,verbose=0,
↳mode='auto')
```

0.4 VGG - 16

0.5 ## VGG Architectures

- **VGG model** was presented as an investigation of the effect of depth on the accuracy of a CNN architecture on a large scale image recognition setting.
- The key idea is to use very small (3 x 3) filters and increasing the number of layers (depth of the network)
- We build and analyse the performance of the VGG - 16 on the MNIST digit dataset.

0.5.1 VGG - 16

- As the name says, it consists of 16 weight layers (13 convolution and 3 Fully connected layers). The original paper used 224 x 224 colored images of the ImageNet dataset (ILSVRC).
- The original model took very less epochs to converge when compared to the other proposals even with a shallow network due to the *implicit regularisation imposed by greater depth and smaller conv filter sizes*.
- The images were randomly cropped and flipped both horizontally & vertically to provide data augmentation.
- We implement VGG-16 on a more simpler MNIST digit dataset.

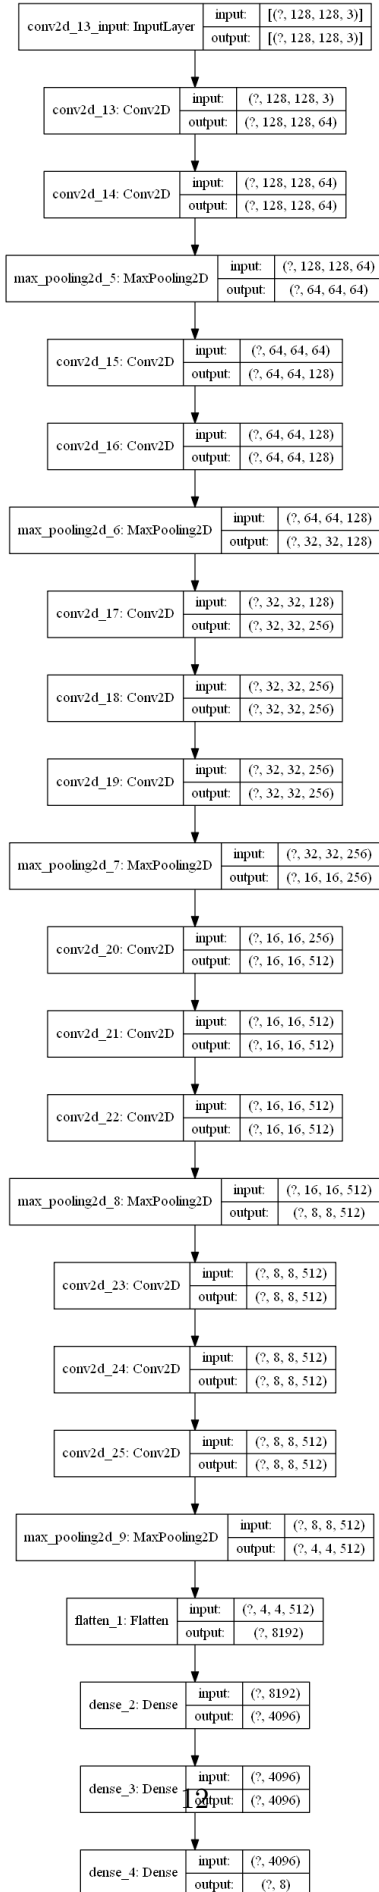
```
[5]: model = Sequential()
model.
    ↳add(Conv2D(input_shape=(n,n,3),filters=64,kernel_size=(3,3),padding="same",
    ↳activation="relu"))
model.add(Conv2D(64,kernel_size=(3,3),padding="same", activation="relu"))
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Conv2D(128, kernel_size=(3,3), padding="same", activation="relu"))
model.add(Conv2D(128, kernel_size=(3,3), padding="same", activation="relu"))
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Conv2D(256, kernel_size=(3,3), padding="same", activation="relu"))
model.add(Conv2D(256, kernel_size=(3,3), padding="same", activation="relu"))
model.add(Conv2D(256, kernel_size=(3,3), padding="same", activation="relu"))
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Conv2D(512, kernel_size=(3,3), padding="same", activation="relu"))
model.add(Conv2D(512, kernel_size=(3,3), padding="same", activation="relu"))
model.add(Conv2D(512, kernel_size=(3,3), padding="same", activation="relu"))
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Conv2D(512, kernel_size=(3,3), padding="same", activation="relu"))
model.add(Conv2D(512, kernel_size=(3,3), padding="same", activation="relu"))
model.add(Conv2D(512, kernel_size=(3,3), padding="same", activation="relu"))
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Flatten())
model.add(Dense(4096,activation="relu"))
model.add(Dense(4096,activation="relu"))
model.add(Dense(units=len(classes), activation="softmax"))
model.summary()
plot_model(model, show_shapes=True, to_file='VGG_16.png')
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_13 (Conv2D)	(None, 128, 128, 64)	1792
conv2d_14 (Conv2D)	(None, 128, 128, 64)	36928
max_pooling2d_5 (MaxPooling2D)	(None, 64, 64, 64)	0
conv2d_15 (Conv2D)	(None, 64, 64, 128)	73856
conv2d_16 (Conv2D)	(None, 64, 64, 128)	147584
max_pooling2d_6 (MaxPooling2D)	(None, 32, 32, 128)	0
conv2d_17 (Conv2D)	(None, 32, 32, 256)	295168
conv2d_18 (Conv2D)	(None, 32, 32, 256)	590080

conv2d_19 (Conv2D)	(None, 32, 32, 256)	590080
max_pooling2d_7 (MaxPooling2D)	(None, 16, 16, 256)	0
conv2d_20 (Conv2D)	(None, 16, 16, 512)	1180160
conv2d_21 (Conv2D)	(None, 16, 16, 512)	2359808
conv2d_22 (Conv2D)	(None, 16, 16, 512)	2359808
max_pooling2d_8 (MaxPooling2D)	(None, 8, 8, 512)	0
conv2d_23 (Conv2D)	(None, 8, 8, 512)	2359808
conv2d_24 (Conv2D)	(None, 8, 8, 512)	2359808
conv2d_25 (Conv2D)	(None, 8, 8, 512)	2359808
max_pooling2d_9 (MaxPooling2D)	(None, 4, 4, 512)	0
flatten_1 (Flatten)	(None, 8192)	0
dense_2 (Dense)	(None, 4096)	33558528
dense_3 (Dense)	(None, 4096)	16781312
dense_4 (Dense)	(None, 8)	32776
=====		
Total params: 65,087,304		
Trainable params: 65,087,304		
Non-trainable params: 0		
=====		

[5]:



```
[12]: model.  
      ↪ compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
```

```
[13]: start_time=t.time()  
      history = model.fit(train_images, train_labels,  
                          validation_data=(test_images, test_labels),  
                          epochs=50, batch_size=64, callbacks=[early])  
      print("\n Training Time: {} seconds".format(t.time()-start_time))  
      model.save('vgg_16_n128_colour.h5')
```

Epoch 1/50

2/157 [...] - ETA: 36s - loss: 2.0795 - accuracy:
0.1094WARNING:tensorflow:Callbacks method `on_train_batch_end` is slow compared
to the batch time (batch time: 0.1237s vs `on_train_batch_end` time: 0.3431s).
Check your callbacks.

157/157 [=====] - 79s 505ms/step - loss: 2.0793 -
accuracy: 0.1195 - val_loss: 2.0783 - val_accuracy: 0.1150

Epoch 2/50

157/157 [=====] - 78s 498ms/step - loss: 2.0742 -
accuracy: 0.1630 - val_loss: 2.0662 - val_accuracy: 0.2225

Epoch 3/50

157/157 [=====] - 79s 501ms/step - loss: 2.0235 -
accuracy: 0.2206 - val_loss: 1.9504 - val_accuracy: 0.2410

Epoch 4/50

157/157 [=====] - 79s 504ms/step - loss: 1.8996 -
accuracy: 0.2752 - val_loss: 1.8147 - val_accuracy: 0.3220

Epoch 5/50

157/157 [=====] - 79s 505ms/step - loss: 1.7712 -
accuracy: 0.3287 - val_loss: 1.8372 - val_accuracy: 0.2985

Epoch 6/50

157/157 [=====] - 79s 505ms/step - loss: 1.6641 -
accuracy: 0.3756 - val_loss: 1.6350 - val_accuracy: 0.3815

Epoch 7/50

157/157 [=====] - 79s 506ms/step - loss: 1.6170 -
accuracy: 0.3905 - val_loss: 1.5708 - val_accuracy: 0.4245

Epoch 8/50

157/157 [=====] - 80s 507ms/step - loss: 1.5750 -
accuracy: 0.4055 - val_loss: 1.5524 - val_accuracy: 0.4270

Epoch 9/50

157/157 [=====] - 80s 509ms/step - loss: 1.5326 -
accuracy: 0.4266 - val_loss: 1.5331 - val_accuracy: 0.4345

Epoch 10/50

157/157 [=====] - 80s 510ms/step - loss: 1.4652 -
accuracy: 0.4490 - val_loss: 1.6354 - val_accuracy: 0.4220

Epoch 11/50

157/157 [=====] - 80s 508ms/step - loss: 1.4322 - accuracy: 0.4704 - val_loss: 1.3916 - val_accuracy: 0.4985
Epoch 12/50
157/157 [=====] - 80s 507ms/step - loss: 1.3614 - accuracy: 0.4945 - val_loss: 1.3791 - val_accuracy: 0.4835
Epoch 13/50
157/157 [=====] - 80s 508ms/step - loss: 1.3120 - accuracy: 0.5141 - val_loss: 1.3238 - val_accuracy: 0.5010
Epoch 14/50
157/157 [=====] - 80s 508ms/step - loss: 1.2695 - accuracy: 0.5234 - val_loss: 1.3821 - val_accuracy: 0.4905
Epoch 15/50
157/157 [=====] - 79s 506ms/step - loss: 1.2385 - accuracy: 0.5355 - val_loss: 1.2715 - val_accuracy: 0.5325
Epoch 16/50
157/157 [=====] - 79s 504ms/step - loss: 1.2113 - accuracy: 0.5472 - val_loss: 1.2334 - val_accuracy: 0.5500
Epoch 17/50
157/157 [=====] - 79s 502ms/step - loss: 1.1892 - accuracy: 0.5495 - val_loss: 1.3401 - val_accuracy: 0.5070
Epoch 18/50
157/157 [=====] - 79s 502ms/step - loss: 1.1713 - accuracy: 0.5628 - val_loss: 1.2190 - val_accuracy: 0.5565
Epoch 19/50
157/157 [=====] - 79s 501ms/step - loss: 1.1435 - accuracy: 0.5706 - val_loss: 1.2083 - val_accuracy: 0.5495
Epoch 20/50
157/157 [=====] - 79s 502ms/step - loss: 1.1250 - accuracy: 0.5773 - val_loss: 1.1920 - val_accuracy: 0.5560
Epoch 21/50
157/157 [=====] - 79s 502ms/step - loss: 1.1200 - accuracy: 0.5831 - val_loss: 1.2096 - val_accuracy: 0.5425
Epoch 22/50
157/157 [=====] - 79s 502ms/step - loss: 1.0897 - accuracy: 0.5893 - val_loss: 1.1601 - val_accuracy: 0.5780
Epoch 23/50
157/157 [=====] - 79s 502ms/step - loss: 1.0803 - accuracy: 0.5961 - val_loss: 1.1696 - val_accuracy: 0.5775
Epoch 24/50
157/157 [=====] - 79s 502ms/step - loss: 1.0554 - accuracy: 0.6085 - val_loss: 1.1435 - val_accuracy: 0.5790
Epoch 25/50
157/157 [=====] - 79s 502ms/step - loss: 1.0447 - accuracy: 0.6078 - val_loss: 1.1547 - val_accuracy: 0.5890
Epoch 26/50
157/157 [=====] - 79s 502ms/step - loss: 1.0314 - accuracy: 0.6166 - val_loss: 1.1319 - val_accuracy: 0.5940
Epoch 27/50

```



157/157 [=====] - 79s 502ms/step - loss: 1.0169 -
accuracy: 0.6199 - val_loss: 1.1631 - val_accuracy: 0.5750
Epoch 28/50
157/157 [=====] - 79s 502ms/step - loss: 0.9966 -
accuracy: 0.6267 - val_loss: 1.1336 - val_accuracy: 0.5955
Epoch 29/50
157/157 [=====] - 79s 502ms/step - loss: 0.9861 -
accuracy: 0.6338 - val_loss: 1.1147 - val_accuracy: 0.5980
Epoch 30/50
157/157 [=====] - 79s 502ms/step - loss: 0.9637 -
accuracy: 0.6434 - val_loss: 1.1129 - val_accuracy: 0.5935
Epoch 31/50
157/157 [=====] - 79s 501ms/step - loss: 0.9383 -
accuracy: 0.6509 - val_loss: 1.1016 - val_accuracy: 0.6070
Epoch 32/50
157/157 [=====] - 79s 501ms/step - loss: 0.9202 -
accuracy: 0.6592 - val_loss: 1.1419 - val_accuracy: 0.6050
Epoch 33/50
157/157 [=====] - 79s 501ms/step - loss: 0.9082 -
accuracy: 0.6604 - val_loss: 1.1485 - val_accuracy: 0.5980
Epoch 34/50
157/157 [=====] - 79s 502ms/step - loss: 0.8884 -
accuracy: 0.6697 - val_loss: 1.1974 - val_accuracy: 0.5750
Epoch 35/50
157/157 [=====] - 79s 501ms/step - loss: 0.8727 -
accuracy: 0.6770 - val_loss: 1.1309 - val_accuracy: 0.6085
Epoch 36/50
157/157 [=====] - 79s 501ms/step - loss: 0.8622 -
accuracy: 0.6781 - val_loss: 1.1134 - val_accuracy: 0.6040
Epoch 37/50
157/157 [=====] - 79s 502ms/step - loss: 0.8357 -
accuracy: 0.6925 - val_loss: 1.1472 - val_accuracy: 0.5990
Epoch 38/50
157/157 [=====] - 79s 502ms/step - loss: 0.8084 -
accuracy: 0.7028 - val_loss: 1.1359 - val_accuracy: 0.6120
Epoch 39/50
157/157 [=====] - 79s 502ms/step - loss: 0.7870 -
accuracy: 0.7057 - val_loss: 1.1493 - val_accuracy: 0.6090
Epoch 40/50
157/157 [=====] - 79s 502ms/step - loss: 0.7797 -
accuracy: 0.7129 - val_loss: 1.1795 - val_accuracy: 0.5895
Epoch 41/50
157/157 [=====] - 79s 501ms/step - loss: 0.7516 -
accuracy: 0.7294 - val_loss: 1.1692 - val_accuracy: 0.6005

```

Training Time: 3267.4202382564545 seconds

```
[14]: score = model.evaluate(test_images, test_labels, verbose=0)
print('Test Loss - ', score[0])
print('Test Accuracy - ', score[1])
graphs(history, score)
```

Test Loss - 1.1692144870758057
Test Accuracy - 0.6004999876022339



Accuracy of the model: 60.05%

0.6 ## VGG - 19

- It is a variation of the VGG-19, with 19 weight layers, 16 from convolution and 3 from the FC

```
[6]: #vgg-19
model2 = Sequential()
model2.
    ↪add(Conv2D(input_shape=(n,n,3),filters=64,kernel_size=(3,3),padding="same",
    ↪activation="relu"))
model2.add(Conv2D(64,kernel_size=(3,3),padding="same", activation="relu"))
model2.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model2.add(Conv2D(128, kernel_size=(3,3), padding="same", activation="relu"))
model2.add(Conv2D(128, kernel_size=(3,3), padding="same", activation="relu"))
model2.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model2.add(Conv2D(256, kernel_size=(3,3), padding="same", activation="relu"))
model2.add(Conv2D(256, kernel_size=(3,3), padding="same", activation="relu"))
model2.add(Conv2D(256, kernel_size=(3,3), padding="same", activation="relu"))
model2.add(Conv2D(256, kernel_size=(3,3), padding="same", activation="relu"))
model2.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model2.add(Conv2D(512, kernel_size=(3,3), padding="same", activation="relu"))
model2.add(Conv2D(512, kernel_size=(3,3), padding="same", activation="relu"))
model2.add(Conv2D(512, kernel_size=(3,3), padding="same", activation="relu"))
model2.add(Conv2D(512, kernel_size=(3,3), padding="same", activation="relu"))
model2.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model2.add(Conv2D(512, kernel_size=(3,3), padding="same", activation="relu"))
model2.add(Conv2D(512, kernel_size=(3,3), padding="same", activation="relu"))
model2.add(Conv2D(512, kernel_size=(3,3), padding="same", activation="relu"))
model2.add(Conv2D(512, kernel_size=(3,3), padding="same", activation="relu"))
model2.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model2.add(Flatten())
model2.add(Dense(4096,activation="relu"))
model2.add(Dense(4096,activation="relu"))
model2.add(Dense(units=len(classes), activation="softmax"))
model2.summary()
plot_model(model2, show_shapes=True, to_file='VGG_19.png')
```

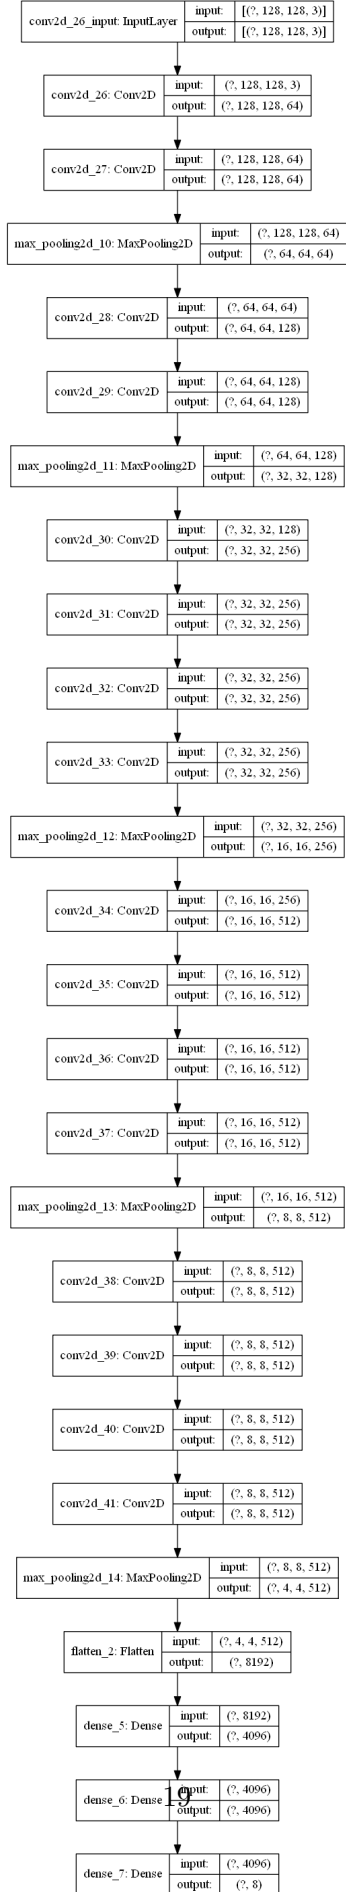
Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_26 (Conv2D)	(None, 128, 128, 64)	1792
conv2d_27 (Conv2D)	(None, 128, 128, 64)	36928
max_pooling2d_10 (MaxPooling)	(None, 64, 64, 64)	0
conv2d_28 (Conv2D)	(None, 64, 64, 128)	73856

conv2d_29 (Conv2D)	(None, 64, 64, 128)	147584
max_pooling2d_11 (MaxPooling)	(None, 32, 32, 128)	0
conv2d_30 (Conv2D)	(None, 32, 32, 256)	295168
conv2d_31 (Conv2D)	(None, 32, 32, 256)	590080
conv2d_32 (Conv2D)	(None, 32, 32, 256)	590080
conv2d_33 (Conv2D)	(None, 32, 32, 256)	590080
max_pooling2d_12 (MaxPooling)	(None, 16, 16, 256)	0
conv2d_34 (Conv2D)	(None, 16, 16, 512)	1180160
conv2d_35 (Conv2D)	(None, 16, 16, 512)	2359808
conv2d_36 (Conv2D)	(None, 16, 16, 512)	2359808
conv2d_37 (Conv2D)	(None, 16, 16, 512)	2359808
max_pooling2d_13 (MaxPooling)	(None, 8, 8, 512)	0
conv2d_38 (Conv2D)	(None, 8, 8, 512)	2359808
conv2d_39 (Conv2D)	(None, 8, 8, 512)	2359808
conv2d_40 (Conv2D)	(None, 8, 8, 512)	2359808
conv2d_41 (Conv2D)	(None, 8, 8, 512)	2359808
max_pooling2d_14 (MaxPooling)	(None, 4, 4, 512)	0
flatten_2 (Flatten)	(None, 8192)	0
dense_5 (Dense)	(None, 4096)	33558528
dense_6 (Dense)	(None, 4096)	16781312
dense_7 (Dense)	(None, 8)	32776

=====
 Total params: 70,397,000
 Trainable params: 70,397,000
 Non-trainable params: 0
 =====

[6]:



```
[15]: model2.  
      ↪ compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
```

```
[16]: start_time=t.time()  
      history2 = model2.fit(train_images, train_labels,  
                           validation_data=(test_images, test_labels),  
                           epochs=100, batch_size=64, callbacks=[early])  
      print("\n Training Time: {} seconds".format(t.time()-start_time))  
      model2.save('vgg_19_n128_colour.h5')
```

Epoch 1/100

2/157 [...] - ETA: 42s - loss: 2.0794 - accuracy:
0.1719WARNING:tensorflow:Callbacks method `on_train_batch_end` is slow compared
to the batch time (batch time: 0.1492s vs `on_train_batch_end` time: 0.4012s).
Check your callbacks.

157/157 [=====] - 93s 593ms/step - loss: 2.0799 -
accuracy: 0.1212 - val_loss: 2.0801 - val_accuracy: 0.1150

Epoch 2/100

157/157 [=====] - 91s 583ms/step - loss: 2.0795 -
accuracy: 0.1250 - val_loss: 2.0800 - val_accuracy: 0.1150

Epoch 3/100

157/157 [=====] - 92s 583ms/step - loss: 2.0792 -
accuracy: 0.1252 - val_loss: 2.0796 - val_accuracy: 0.1150

Epoch 4/100

157/157 [=====] - 92s 583ms/step - loss: 2.0788 -
accuracy: 0.1253 - val_loss: 2.0792 - val_accuracy: 0.1150

Epoch 5/100

157/157 [=====] - 92s 584ms/step - loss: 2.0781 -
accuracy: 0.1378 - val_loss: 2.0783 - val_accuracy: 0.1430

Epoch 6/100

157/157 [=====] - 92s 584ms/step - loss: 2.0765 -
accuracy: 0.1804 - val_loss: 2.0760 - val_accuracy: 0.1695

Epoch 7/100

157/157 [=====] - 92s 583ms/step - loss: 2.0733 -
accuracy: 0.2111 - val_loss: 2.0715 - val_accuracy: 0.2265

Epoch 8/100

157/157 [=====] - 92s 584ms/step - loss: 2.0672 -
accuracy: 0.2286 - val_loss: 2.0633 - val_accuracy: 0.2275

Epoch 9/100

157/157 [=====] - 92s 584ms/step - loss: 2.0567 -
accuracy: 0.2237 - val_loss: 2.0496 - val_accuracy: 0.2185

Epoch 10/100

157/157 [=====] - 92s 584ms/step - loss: 2.0334 -
accuracy: 0.2161 - val_loss: 2.0054 - val_accuracy: 0.2190

Epoch 11/100

157/157 [=====] - 92s 584ms/step - loss: 1.9654 - accuracy: 0.2113 - val_loss: 1.9493 - val_accuracy: 0.2145
Epoch 12/100
157/157 [=====] - 92s 584ms/step - loss: 1.9196 - accuracy: 0.2216 - val_loss: 1.9099 - val_accuracy: 0.2320
Epoch 13/100
157/157 [=====] - 92s 584ms/step - loss: 1.8858 - accuracy: 0.2329 - val_loss: 1.8826 - val_accuracy: 0.2645
Epoch 14/100
157/157 [=====] - 92s 583ms/step - loss: 1.8578 - accuracy: 0.2500 - val_loss: 1.8792 - val_accuracy: 0.2585
Epoch 15/100
157/157 [=====] - 92s 583ms/step - loss: 1.8396 - accuracy: 0.2637 - val_loss: 1.8417 - val_accuracy: 0.2335
Epoch 16/100
157/157 [=====] - 92s 583ms/step - loss: 1.7945 - accuracy: 0.2843 - val_loss: 1.7679 - val_accuracy: 0.3005
Epoch 17/100
157/157 [=====] - 92s 583ms/step - loss: 1.7651 - accuracy: 0.3024 - val_loss: 1.7336 - val_accuracy: 0.3135
Epoch 18/100
157/157 [=====] - 91s 583ms/step - loss: 1.7634 - accuracy: 0.3059 - val_loss: 1.7508 - val_accuracy: 0.3005
Epoch 19/100
157/157 [=====] - 91s 582ms/step - loss: 1.7313 - accuracy: 0.3252 - val_loss: 1.7216 - val_accuracy: 0.3285
Epoch 20/100
157/157 [=====] - 92s 583ms/step - loss: 1.6920 - accuracy: 0.3571 - val_loss: 1.6626 - val_accuracy: 0.3705
Epoch 21/100
157/157 [=====] - 91s 583ms/step - loss: 1.6443 - accuracy: 0.3806 - val_loss: 1.6287 - val_accuracy: 0.4015
Epoch 22/100
157/157 [=====] - 91s 583ms/step - loss: 1.6009 - accuracy: 0.3885 - val_loss: 1.6152 - val_accuracy: 0.4050
Epoch 23/100
157/157 [=====] - 92s 583ms/step - loss: 1.5741 - accuracy: 0.4068 - val_loss: 1.5335 - val_accuracy: 0.4585
Epoch 24/100
157/157 [=====] - 92s 584ms/step - loss: 1.5687 - accuracy: 0.4103 - val_loss: 1.6175 - val_accuracy: 0.3980
Epoch 25/100
157/157 [=====] - 92s 584ms/step - loss: 1.5411 - accuracy: 0.4253 - val_loss: 1.5061 - val_accuracy: 0.4530
Epoch 26/100
157/157 [=====] - 92s 584ms/step - loss: 1.5144 - accuracy: 0.4333 - val_loss: 1.5174 - val_accuracy: 0.4345
Epoch 27/100

157/157 [=====] - 92s 584ms/step - loss: 1.4827 -
accuracy: 0.4424 - val_loss: 1.4914 - val_accuracy: 0.4400
Epoch 28/100
157/157 [=====] - 92s 584ms/step - loss: 1.4807 -
accuracy: 0.4412 - val_loss: 1.4906 - val_accuracy: 0.4545
Epoch 29/100
157/157 [=====] - 92s 584ms/step - loss: 1.4696 -
accuracy: 0.4435 - val_loss: 1.4602 - val_accuracy: 0.4490
Epoch 30/100
157/157 [=====] - 92s 584ms/step - loss: 1.4568 -
accuracy: 0.4506 - val_loss: 1.4696 - val_accuracy: 0.4545
Epoch 31/100
157/157 [=====] - 92s 584ms/step - loss: 1.4229 -
accuracy: 0.4599 - val_loss: 1.4777 - val_accuracy: 0.4465
Epoch 32/100
157/157 [=====] - 92s 584ms/step - loss: 1.4187 -
accuracy: 0.4634 - val_loss: 1.3999 - val_accuracy: 0.4795
Epoch 33/100
157/157 [=====] - 92s 584ms/step - loss: 1.3875 -
accuracy: 0.4725 - val_loss: 1.5598 - val_accuracy: 0.4110
Epoch 34/100
157/157 [=====] - 92s 585ms/step - loss: 1.3933 -
accuracy: 0.4722 - val_loss: 1.4840 - val_accuracy: 0.4695
Epoch 35/100
157/157 [=====] - 92s 585ms/step - loss: 1.3894 -
accuracy: 0.4736 - val_loss: 1.4072 - val_accuracy: 0.4750
Epoch 36/100
157/157 [=====] - 92s 585ms/step - loss: 1.3527 -
accuracy: 0.4898 - val_loss: 1.3661 - val_accuracy: 0.4940
Epoch 37/100
157/157 [=====] - 92s 584ms/step - loss: 1.3298 -
accuracy: 0.4947 - val_loss: 1.5163 - val_accuracy: 0.4385
Epoch 38/100
157/157 [=====] - 92s 584ms/step - loss: 1.3294 -
accuracy: 0.4948 - val_loss: 1.3582 - val_accuracy: 0.5070
Epoch 39/100
157/157 [=====] - 92s 584ms/step - loss: 1.3097 -
accuracy: 0.5100 - val_loss: 1.3224 - val_accuracy: 0.5140
Epoch 40/100
157/157 [=====] - 92s 584ms/step - loss: 1.2796 -
accuracy: 0.5188 - val_loss: 1.3038 - val_accuracy: 0.5240
Epoch 41/100
157/157 [=====] - 92s 584ms/step - loss: 1.2630 -
accuracy: 0.5206 - val_loss: 1.3179 - val_accuracy: 0.5125
Epoch 42/100
157/157 [=====] - 92s 584ms/step - loss: 1.2755 -
accuracy: 0.5162 - val_loss: 1.3314 - val_accuracy: 0.5070
Epoch 43/100

157/157 [=====] - 92s 584ms/step - loss: 1.2420 -
accuracy: 0.5343 - val_loss: 1.3027 - val_accuracy: 0.5135
Epoch 44/100
157/157 [=====] - 92s 585ms/step - loss: 1.2390 -
accuracy: 0.5344 - val_loss: 1.3027 - val_accuracy: 0.5205
Epoch 45/100
157/157 [=====] - 92s 586ms/step - loss: 1.2264 -
accuracy: 0.5393 - val_loss: 1.2874 - val_accuracy: 0.5265
Epoch 46/100
157/157 [=====] - 92s 585ms/step - loss: 1.2034 -
accuracy: 0.5450 - val_loss: 1.2806 - val_accuracy: 0.5295
Epoch 47/100
157/157 [=====] - 92s 585ms/step - loss: 1.1931 -
accuracy: 0.5552 - val_loss: 1.2502 - val_accuracy: 0.5460
Epoch 48/100
157/157 [=====] - 92s 585ms/step - loss: 1.1922 -
accuracy: 0.5563 - val_loss: 1.2473 - val_accuracy: 0.5485
Epoch 49/100
157/157 [=====] - 92s 585ms/step - loss: 1.1754 -
accuracy: 0.5643 - val_loss: 1.2344 - val_accuracy: 0.5615
Epoch 50/100
157/157 [=====] - 92s 585ms/step - loss: 1.1658 -
accuracy: 0.5669 - val_loss: 1.2634 - val_accuracy: 0.5370
Epoch 51/100
157/157 [=====] - 92s 585ms/step - loss: 1.1494 -
accuracy: 0.5706 - val_loss: 1.3133 - val_accuracy: 0.5160
Epoch 52/100
157/157 [=====] - 92s 585ms/step - loss: 1.1530 -
accuracy: 0.5703 - val_loss: 1.2313 - val_accuracy: 0.5520
Epoch 53/100
157/157 [=====] - 92s 585ms/step - loss: 1.1325 -
accuracy: 0.5771 - val_loss: 1.2257 - val_accuracy: 0.5565
Epoch 54/100
157/157 [=====] - 92s 584ms/step - loss: 1.1324 -
accuracy: 0.5780 - val_loss: 1.2254 - val_accuracy: 0.5605
Epoch 55/100
157/157 [=====] - 92s 584ms/step - loss: 1.1285 -
accuracy: 0.5794 - val_loss: 1.2001 - val_accuracy: 0.5635
Epoch 56/100
157/157 [=====] - 92s 584ms/step - loss: 1.1303 -
accuracy: 0.5774 - val_loss: 1.2132 - val_accuracy: 0.5575
Epoch 57/100
157/157 [=====] - 92s 584ms/step - loss: 1.1058 -
accuracy: 0.5855 - val_loss: 1.2355 - val_accuracy: 0.5465
Epoch 58/100
157/157 [=====] - 92s 584ms/step - loss: 1.1100 -
accuracy: 0.5875 - val_loss: 1.2463 - val_accuracy: 0.5425
Epoch 59/100

157/157 [=====] - 92s 583ms/step - loss: 1.0879 -
accuracy: 0.5904 - val_loss: 1.1773 - val_accuracy: 0.5780
Epoch 60/100
157/157 [=====] - 92s 583ms/step - loss: 1.0811 -
accuracy: 0.5979 - val_loss: 1.2008 - val_accuracy: 0.5670
Epoch 61/100
157/157 [=====] - 92s 583ms/step - loss: 1.0778 -
accuracy: 0.6013 - val_loss: 1.1995 - val_accuracy: 0.5675
Epoch 62/100
157/157 [=====] - 92s 583ms/step - loss: 1.0604 -
accuracy: 0.6075 - val_loss: 1.1730 - val_accuracy: 0.5800
Epoch 63/100
157/157 [=====] - 92s 584ms/step - loss: 1.0628 -
accuracy: 0.6061 - val_loss: 1.1739 - val_accuracy: 0.5815
Epoch 64/100
157/157 [=====] - 92s 583ms/step - loss: 1.0557 -
accuracy: 0.6084 - val_loss: 1.1742 - val_accuracy: 0.5925
Epoch 65/100
157/157 [=====] - 91s 582ms/step - loss: 1.0433 -
accuracy: 0.6159 - val_loss: 1.1811 - val_accuracy: 0.5775
Epoch 66/100
157/157 [=====] - 91s 582ms/step - loss: 1.0370 -
accuracy: 0.6159 - val_loss: 1.1909 - val_accuracy: 0.5730
Epoch 67/100
157/157 [=====] - 91s 582ms/step - loss: 1.0319 -
accuracy: 0.6183 - val_loss: 1.2389 - val_accuracy: 0.5430
Epoch 68/100
157/157 [=====] - 91s 582ms/step - loss: 1.0275 -
accuracy: 0.6174 - val_loss: 1.1599 - val_accuracy: 0.5890
Epoch 69/100
157/157 [=====] - 92s 583ms/step - loss: 1.0083 -
accuracy: 0.6276 - val_loss: 1.1626 - val_accuracy: 0.5785
Epoch 70/100
157/157 [=====] - 92s 583ms/step - loss: 0.9981 -
accuracy: 0.6323 - val_loss: 1.1769 - val_accuracy: 0.5720
Epoch 71/100
157/157 [=====] - 92s 584ms/step - loss: 1.0001 -
accuracy: 0.6298 - val_loss: 1.1490 - val_accuracy: 0.5890
Epoch 72/100
157/157 [=====] - 92s 585ms/step - loss: 0.9848 -
accuracy: 0.6363 - val_loss: 1.1828 - val_accuracy: 0.5740
Epoch 73/100
157/157 [=====] - 92s 584ms/step - loss: 0.9931 -
accuracy: 0.6310 - val_loss: 1.1895 - val_accuracy: 0.5740
Epoch 74/100
157/157 [=====] - 92s 584ms/step - loss: 0.9691 -
accuracy: 0.6392 - val_loss: 1.2396 - val_accuracy: 0.5505
Epoch 75/100


```

157/157 [=====] - 92s 585ms/step - loss: 0.9655 -
accuracy: 0.6401 - val_loss: 1.1777 - val_accuracy: 0.5950
Epoch 76/100
157/157 [=====] - 92s 585ms/step - loss: 0.9684 -
accuracy: 0.6362 - val_loss: 1.1879 - val_accuracy: 0.5785
Epoch 77/100
157/157 [=====] - 92s 585ms/step - loss: 0.9606 -
accuracy: 0.6436 - val_loss: 1.1916 - val_accuracy: 0.5840
Epoch 78/100
157/157 [=====] - 92s 585ms/step - loss: 0.9581 -
accuracy: 0.6420 - val_loss: 1.1564 - val_accuracy: 0.5985
Epoch 79/100
157/157 [=====] - 92s 585ms/step - loss: 0.9392 -
accuracy: 0.6570 - val_loss: 1.1907 - val_accuracy: 0.5835
Epoch 80/100
157/157 [=====] - 92s 585ms/step - loss: 0.9115 -
accuracy: 0.6632 - val_loss: 1.2120 - val_accuracy: 0.5850
Epoch 81/100
157/157 [=====] - 92s 585ms/step - loss: 0.9215 -
accuracy: 0.6578 - val_loss: 1.2593 - val_accuracy: 0.5580

```

Training Time: 7479.614980459213 seconds


```


[17]: score2 = model2.evaluate(test_images,test_labels,verbose=0)
print('Test Loss - ',score2[0])
print('Test Accuracy - ',score2[1])
graphs(history2,score2)

```

Test Loss - 1.259333610534668

Test Accuracy - 0.5580000281333923





Accuracy of the model: 55.80%

0.7 ## AlexNet

- It is comparatively a very small network with only 5 convolution layers and varying filter sizes.

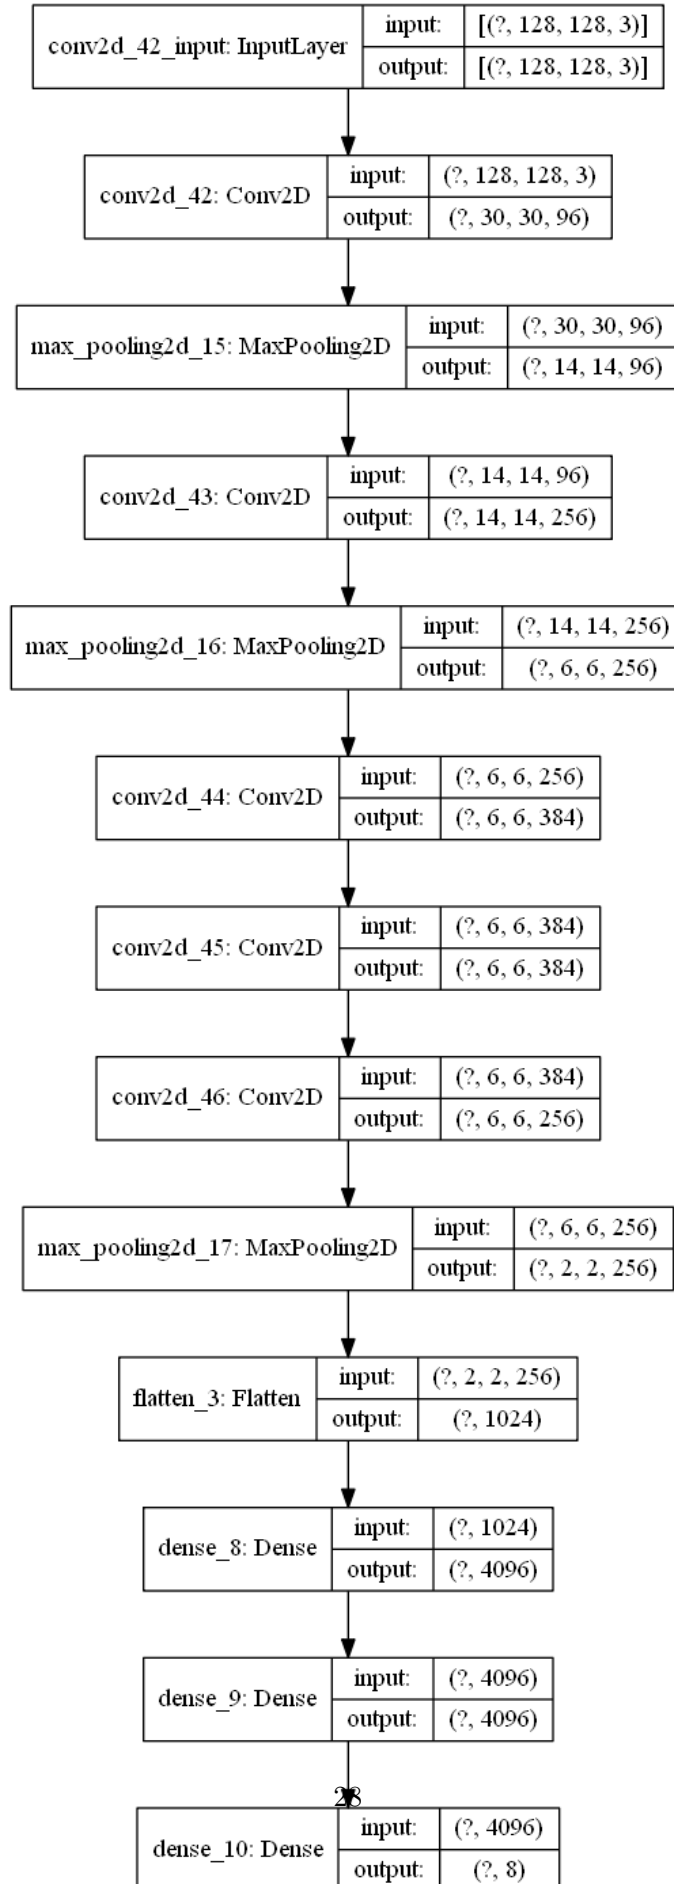
```
[7]: #alexnet
model3 = Sequential()
model3.add(Conv2D(input_shape=(n,n,3),filters=96,kernel_size=(11,11),strides=(4,4),padding="valid",
    ↪activation="relu"))
model3.add(MaxPooling2D(pool_size=(3,3),strides=(2,2)))
model3.add(Conv2D(256, kernel_size=(5,5), padding="same",strides=(1,1),
    ↪activation="relu"))
model3.add(MaxPooling2D(pool_size=(3,3),strides=(2,2)))
model3.add(Conv2D(384, kernel_size=(3,3), padding="same", strides=(1,1),
    ↪activation="relu"))
model3.add(Conv2D(384, kernel_size=(3,3), padding="same", strides=(1,1),
    ↪activation="relu"))
model3.add(Conv2D(256, kernel_size=(3,3), padding="same", strides=(1,1),
    ↪activation="relu"))
model3.add(MaxPooling2D(pool_size=(3,3),strides=(2,2)))
model3.add(Flatten())
model3.add(Dense(4096,activation="relu"))
model3.add(Dense(4096,activation="relu"))
model3.add(Dense(units=len(classes), activation="softmax"))
```

```
model3.summary()
plot_model(model3, show_shapes=True, to_file='ALEX_NET.png')
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
conv2d_42 (Conv2D)	(None, 30, 30, 96)	34944
max_pooling2d_15 (MaxPooling)	(None, 14, 14, 96)	0
conv2d_43 (Conv2D)	(None, 14, 14, 256)	614656
max_pooling2d_16 (MaxPooling)	(None, 6, 6, 256)	0
conv2d_44 (Conv2D)	(None, 6, 6, 384)	885120
conv2d_45 (Conv2D)	(None, 6, 6, 384)	1327488
conv2d_46 (Conv2D)	(None, 6, 6, 256)	884992
max_pooling2d_17 (MaxPooling)	(None, 2, 2, 256)	0
flatten_3 (Flatten)	(None, 1024)	0
dense_8 (Dense)	(None, 4096)	4198400
dense_9 (Dense)	(None, 4096)	16781312
dense_10 (Dense)	(None, 8)	32776
Total params: 24,759,688		
Trainable params: 24,759,688		
Non-trainable params: 0		

[7]:



```
[12]: model3.  
      ↪ compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
```

```
[13]: start_time=t.time()  
      history3 = model3.fit(train_images, train_labels,  
                           validation_data=(test_images, test_labels),  
                           epochs=100, batch_size=64, callbacks=[early])  
      print("\n Training Time: {} seconds".format(t.time()-start_time))  
      model3.save('alexnet_n128_colour.h5')
```

Epoch 1/100

157/157 [=====] - 8s 54ms/step - loss: 1.9990 -
accuracy: 0.2024 - val_loss: 1.8580 - val_accuracy: 0.3075

Epoch 2/100

157/157 [=====] - 8s 49ms/step - loss: 1.7447 -
accuracy: 0.3254 - val_loss: 1.9200 - val_accuracy: 0.2960

Epoch 3/100

157/157 [=====] - 8s 49ms/step - loss: 1.5547 -
accuracy: 0.4045 - val_loss: 1.4474 - val_accuracy: 0.4690

Epoch 4/100

157/157 [=====] - 8s 49ms/step - loss: 1.3725 -
accuracy: 0.4887 - val_loss: 1.3616 - val_accuracy: 0.4940

Epoch 5/100

157/157 [=====] - 8s 49ms/step - loss: 1.2997 -
accuracy: 0.5190 - val_loss: 1.4937 - val_accuracy: 0.4405

Epoch 6/100

157/157 [=====] - 8s 49ms/step - loss: 1.2072 -
accuracy: 0.5506 - val_loss: 1.2207 - val_accuracy: 0.5485

Epoch 7/100

157/157 [=====] - 8s 49ms/step - loss: 1.1483 -
accuracy: 0.5744 - val_loss: 1.2489 - val_accuracy: 0.5485

Epoch 8/100

157/157 [=====] - 8s 49ms/step - loss: 1.1211 -
accuracy: 0.5842 - val_loss: 1.2298 - val_accuracy: 0.5440

Epoch 9/100

157/157 [=====] - 8s 49ms/step - loss: 1.0981 -
accuracy: 0.5937 - val_loss: 1.1190 - val_accuracy: 0.5950

Epoch 10/100

157/157 [=====] - 8s 49ms/step - loss: 1.0569 -
accuracy: 0.6077 - val_loss: 1.0918 - val_accuracy: 0.5980

Epoch 11/100

157/157 [=====] - 8s 49ms/step - loss: 1.0337 -
accuracy: 0.6184 - val_loss: 1.0918 - val_accuracy: 0.6025

Epoch 12/100

157/157 [=====] - 8s 49ms/step - loss: 1.0029 -

accuracy: 0.6268 - val_loss: 1.0568 - val_accuracy: 0.6110
 Epoch 13/100
 157/157 [=====] - 8s 49ms/step - loss: 0.9875 -
 accuracy: 0.6340 - val_loss: 1.1050 - val_accuracy: 0.5945
 Epoch 14/100
 157/157 [=====] - 8s 50ms/step - loss: 0.9688 -
 accuracy: 0.6386 - val_loss: 1.0247 - val_accuracy: 0.6300
 Epoch 15/100
 157/157 [=====] - 8s 50ms/step - loss: 0.9435 -
 accuracy: 0.6514 - val_loss: 1.0384 - val_accuracy: 0.6170
 Epoch 16/100
 157/157 [=====] - 8s 50ms/step - loss: 0.9284 -
 accuracy: 0.6562 - val_loss: 1.0282 - val_accuracy: 0.6310
 Epoch 17/100
 157/157 [=====] - 8s 50ms/step - loss: 0.9095 -
 accuracy: 0.6686 - val_loss: 1.0407 - val_accuracy: 0.6135
 Epoch 18/100
 157/157 [=====] - 8s 50ms/step - loss: 0.9018 -
 accuracy: 0.6666 - val_loss: 0.9957 - val_accuracy: 0.6380
 Epoch 19/100
 157/157 [=====] - 8s 50ms/step - loss: 0.8751 -
 accuracy: 0.6814 - val_loss: 1.0291 - val_accuracy: 0.6225
 Epoch 20/100
 157/157 [=====] - 8s 50ms/step - loss: 0.8671 -
 accuracy: 0.6882 - val_loss: 1.0147 - val_accuracy: 0.6305
 Epoch 21/100
 157/157 [=====] - 8s 50ms/step - loss: 0.8510 -
 accuracy: 0.6868 - val_loss: 0.9729 - val_accuracy: 0.6535
 Epoch 22/100
 157/157 [=====] - 8s 50ms/step - loss: 0.8365 -
 accuracy: 0.6929 - val_loss: 0.9775 - val_accuracy: 0.6420
 Epoch 23/100
 157/157 [=====] - 8s 50ms/step - loss: 0.8265 -
 accuracy: 0.6949 - val_loss: 0.9778 - val_accuracy: 0.6490
 Epoch 24/100
 157/157 [=====] - 8s 50ms/step - loss: 0.8043 -
 accuracy: 0.7074 - val_loss: 0.9656 - val_accuracy: 0.6450
 Epoch 25/100
 157/157 [=====] - 8s 50ms/step - loss: 0.8037 -
 accuracy: 0.7039 - val_loss: 0.9902 - val_accuracy: 0.6350
 Epoch 26/100
 157/157 [=====] - 8s 50ms/step - loss: 0.7811 -
 accuracy: 0.7172 - val_loss: 0.9436 - val_accuracy: 0.6605
 Epoch 27/100
 157/157 [=====] - 8s 50ms/step - loss: 0.7828 -
 accuracy: 0.7127 - val_loss: 0.9693 - val_accuracy: 0.6560
 Epoch 28/100
 157/157 [=====] - 8s 50ms/step - loss: 0.7736 -

accuracy: 0.7191 - val_loss: 0.9208 - val_accuracy: 0.6695
 Epoch 29/100
 157/157 [=====] - 8s 50ms/step - loss: 0.7562 -
 accuracy: 0.7235 - val_loss: 0.9628 - val_accuracy: 0.6490
 Epoch 30/100
 157/157 [=====] - 8s 50ms/step - loss: 0.7540 -
 accuracy: 0.7248 - val_loss: 0.9729 - val_accuracy: 0.6535
 Epoch 31/100
 157/157 [=====] - 8s 50ms/step - loss: 0.7318 -
 accuracy: 0.7329 - val_loss: 0.9215 - val_accuracy: 0.6710
 Epoch 32/100
 157/157 [=====] - 8s 50ms/step - loss: 0.7240 -
 accuracy: 0.7362 - val_loss: 0.9316 - val_accuracy: 0.6540
 Epoch 33/100
 157/157 [=====] - 8s 50ms/step - loss: 0.7168 -
 accuracy: 0.7379 - val_loss: 0.9241 - val_accuracy: 0.6630
 Epoch 34/100
 157/157 [=====] - 8s 50ms/step - loss: 0.7003 -
 accuracy: 0.7461 - val_loss: 0.9128 - val_accuracy: 0.6760
 Epoch 35/100
 157/157 [=====] - 8s 50ms/step - loss: 0.6994 -
 accuracy: 0.7480 - val_loss: 0.9275 - val_accuracy: 0.6625
 Epoch 36/100
 157/157 [=====] - 8s 50ms/step - loss: 0.6839 -
 accuracy: 0.7498 - val_loss: 0.9375 - val_accuracy: 0.6605
 Epoch 37/100
 157/157 [=====] - 8s 50ms/step - loss: 0.6914 -
 accuracy: 0.7466 - val_loss: 0.8952 - val_accuracy: 0.6745
 Epoch 38/100
 157/157 [=====] - 8s 50ms/step - loss: 0.6652 -
 accuracy: 0.7586 - val_loss: 0.9353 - val_accuracy: 0.6680
 Epoch 39/100
 157/157 [=====] - 8s 50ms/step - loss: 0.6599 -
 accuracy: 0.7596 - val_loss: 0.8892 - val_accuracy: 0.6780
 Epoch 40/100
 157/157 [=====] - 8s 50ms/step - loss: 0.6473 -
 accuracy: 0.7672 - val_loss: 0.9025 - val_accuracy: 0.6755
 Epoch 41/100
 157/157 [=====] - 8s 50ms/step - loss: 0.6494 -
 accuracy: 0.7653 - val_loss: 0.9057 - val_accuracy: 0.6815
 Epoch 42/100
 157/157 [=====] - 8s 50ms/step - loss: 0.6291 -
 accuracy: 0.7718 - val_loss: 0.9044 - val_accuracy: 0.6760
 Epoch 43/100
 157/157 [=====] - 8s 50ms/step - loss: 0.6266 -
 accuracy: 0.7715 - val_loss: 0.8987 - val_accuracy: 0.6815
 Epoch 44/100
 157/157 [=====] - 8s 50ms/step - loss: 0.6158 -

```

accuracy: 0.7758 - val_loss: 0.8964 - val_accuracy: 0.6835
Epoch 45/100
157/157 [=====] - 8s 50ms/step - loss: 0.6119 -
accuracy: 0.7774 - val_loss: 0.9166 - val_accuracy: 0.6785
Epoch 46/100
157/157 [=====] - 8s 50ms/step - loss: 0.6160 -
accuracy: 0.7764 - val_loss: 0.8826 - val_accuracy: 0.6865
Epoch 47/100
157/157 [=====] - 8s 50ms/step - loss: 0.6048 -
accuracy: 0.7832 - val_loss: 0.9630 - val_accuracy: 0.6570
Epoch 48/100
157/157 [=====] - 8s 50ms/step - loss: 0.5962 -
accuracy: 0.7837 - val_loss: 0.9109 - val_accuracy: 0.6780
Epoch 49/100
157/157 [=====] - 8s 50ms/step - loss: 0.5893 -
accuracy: 0.7881 - val_loss: 0.8983 - val_accuracy: 0.6795
Epoch 50/100
157/157 [=====] - 8s 50ms/step - loss: 0.5733 -
accuracy: 0.7960 - val_loss: 0.9185 - val_accuracy: 0.6675
Epoch 51/100
157/157 [=====] - 8s 50ms/step - loss: 0.5628 -
accuracy: 0.7992 - val_loss: 0.9101 - val_accuracy: 0.6795
Epoch 52/100
157/157 [=====] - 8s 50ms/step - loss: 0.5677 -
accuracy: 0.7937 - val_loss: 0.9100 - val_accuracy: 0.6810
Epoch 53/100
157/157 [=====] - 8s 50ms/step - loss: 0.5553 -
accuracy: 0.7989 - val_loss: 0.9205 - val_accuracy: 0.6795
Epoch 54/100
157/157 [=====] - 8s 50ms/step - loss: 0.5495 -
accuracy: 0.8040 - val_loss: 0.9032 - val_accuracy: 0.6765
Epoch 55/100
157/157 [=====] - 8s 50ms/step - loss: 0.5504 -
accuracy: 0.8031 - val_loss: 0.8839 - val_accuracy: 0.6905- loss: 0.5526 - accu
Epoch 56/100
157/157 [=====] - 8s 50ms/step - loss: 0.5406 -
accuracy: 0.8061 - val_loss: 0.8885 - val_accuracy: 0.6850

```

Training Time: 447.03455662727356 seconds

```



[14]: score3 = model3.evaluate(test_images,test_labels,verbose=0)
print('Test Loss - ',score3[0])
print('Test Accuracy - ',score3[1])
graphs(history3,score3)

```

```

Test Loss - 0.888465404510498
Test Accuracy - 0.6850000023841858

```

Accuracy of the model: 68.50%

0.8 Inception Module

- Some ideas on the inception module implementation from: [1]
F. SHAIKH, ``Inception Network | Implementation Of GoogleNet In Keras'', Analytics Vidhya, 2020. [Online]. Available: <https://www.analyticsvidhya.com/blog/2018/10/understanding-inception-network-from-scratch/> [Accessed: 12- Aug- 2020]

```
[8]: # function for creating a projected inception module (performance improved)
def inception_module(layer_in, f1, f2_in, f2_out, f3_in, f3_out, f4_out):
    conv1 = Conv2D(f1, (1,1), padding='same', activation='relu')(layer_in)
    conv3 = Conv2D(f2_in, (1,1), padding='same', activation='relu')(layer_in)
    conv3 = Conv2D(f2_out, (3,3), padding='same', activation='relu')(conv3)
    conv5 = Conv2D(f3_in, (1,1), padding='same', activation='relu')(layer_in)
    conv5 = Conv2D(f3_out, (5,5), padding='same', activation='relu')(conv5)
    pool = MaxPooling2D((3,3), strides=(1,1), padding='same')(layer_in)
    pool = Conv2D(f4_out, (1,1), padding='same', activation='relu')(pool)
    # concatenate filters, assumes filters/channels last
    layer_out = concatenate([conv1, conv3, conv5, pool], axis=-1)
    return layer_out
```

0.9 Let's test our function

```
[9]: visible = Input(shape=(256, 256, 3))
# add inception block 1
layer = inception_module(visible, 64, 96, 128, 16, 32, 32)
# create model
model = Model(inputs=visible, outputs=layer)
# summarize model
model.summary()
# plot model architecture
plot_model(model, show_shapes=True)
```

Model: "functional_1"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 256, 256, 3)]	0	
conv2d_1 (Conv2D)	(None, 256, 256, 96)	384	input_1[0][0]
conv2d_3 (Conv2D)	(None, 256, 256, 16)	64	input_1[0][0]
max_pooling2d (MaxPooling2D)	(None, 256, 256, 3)	0	input_1[0][0]
conv2d (Conv2D)	(None, 256, 256, 64)	256	input_1[0][0]
conv2d_2 (Conv2D)	(None, 256, 256, 128)	110720	conv2d_1[0][0]

```

-----
conv2d_4 (Conv2D)                (None, 256, 256, 32) 12832      conv2d_3[0][0]
-----
conv2d_5 (Conv2D)                (None, 256, 256, 32) 128
max_pooling2d[0][0]
-----
concatenate (Concatenate)        (None, 256, 256, 256) 0      conv2d[0][0]
                                   conv2d_2[0][0]
                                   conv2d_4[0][0]
                                   conv2d_5[0][0]
=====
=====

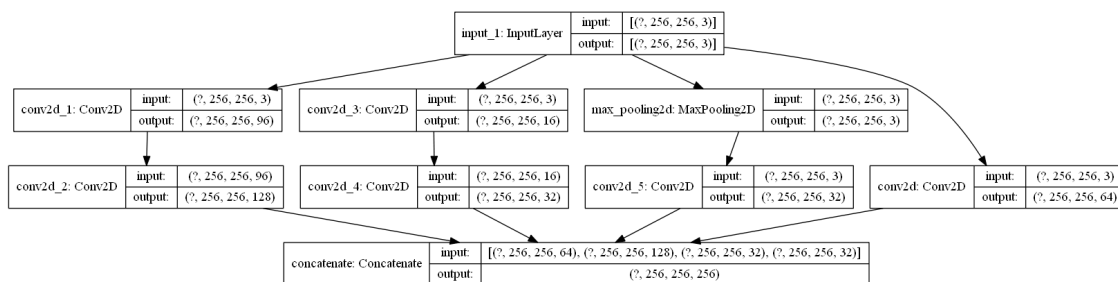
```

```

Total params: 124,384
Trainable params: 124,384
Non-trainable params: 0
-----
-----

```

[9]:



0.10 GoogleNet using inception module

0.10.1 Relevant sources

[1] C. Szegedy et al., ``Going deeper with convolutions,`` 2015 IEEE Conference on Computer Vision and Pattern Recognition (CVPR), Boston, MA, 2015, pp. 1-9, doi: 10.1109/CVPR.2015.7298594.

```

[10]: kernel_init = keras.initializers.glorot_uniform()
      bias_init = keras.initializers.Constant(value=0.2)

```

```

[12]: # define model input
      input_layer = Input(shape=(128, 128, 3))

```

```

x = Conv2D(64, (7, 7), padding='same', strides=(2, 2),
    ↪activation='relu',kernel_initializer=kernel_init,
    ↪bias_initializer=bias_init)(input_layer)
x = MaxPool2D((3, 3), padding='same', strides=(2, 2))(x)
x = Conv2D(64, (1, 1), padding='same', strides=(1, 1), activation='relu')(x)
x = Conv2D(192, (3, 3), padding='same', strides=(1, 1), activation='relu')(x)
x = MaxPool2D((3, 3), padding='same', strides=(2, 2))(x)

x = inception_module(x, 64, 96, 128, 16, 32, 32) #3a
x = inception_module(x, 128, 128, 192, 32, 96, 64) #3b
x = MaxPool2D((3, 3), padding='same', strides=(2, 2))(x)

x = inception_module(x, 192, 96, 208, 16, 48, 64) #4a

x1 = AveragePooling2D((5, 5), strides=3)(x)
x1 = Conv2D(128, (1, 1), padding='same', activation='relu')(x1)
x1 = Flatten()(x1)
x1 = Dense(1024, activation='relu')(x1)
x1 = Dropout(0.7)(x1)
x1 = Dense(len(classes), activation='softmax', name="Softmax_1")(x1)

x = inception_module(x, 160, 112, 224, 24, 64, 64) #4b
x = inception_module(x, 128, 128, 256, 24, 64, 64) #4c
x = inception_module(x, 112, 144, 288, 32, 64, 64) #4d

x2 = AveragePooling2D((5, 5), strides=3)(x)
x2 = Conv2D(128, (1, 1), padding='same', activation='relu')(x2)
x2 = Flatten()(x2)
x2 = Dense(1024, activation='relu')(x2)
x2 = Dropout(0.7)(x2)
x2 = Dense(len(classes), activation='softmax',name="Softmax_2")(x2)

x = inception_module(x, 256, 160, 320, 32, 128, 128) #4e

x = MaxPool2D((3, 3), padding='same', strides=(2, 2))(x)

x = inception_module(x, 256, 160, 320, 32, 128, 128) #5a
x = inception_module(x, 384, 192, 384, 48, 128, 128 ) #5b

x = GlobalAveragePooling2D()(x)
x = Dropout(0.4)(x)
x = Dense(len(classes), activation='softmax',name="OUTPUT")(x)

model = Model(input_layer, [x, x1, x2], name='inception_v1')
model.summary()
plot_model(model)

```

Model: "inception_v1"

Layer (type)	Output Shape	Param #	Connected to
input_2 (InputLayer)	[(None, 128, 128, 3)]	0	
conv2d_6 (Conv2D)	(None, 64, 64, 64)	9472	input_2[0][0]
max_pooling2d_1 (MaxPooling2D)	(None, 32, 32, 64)	0	conv2d_6[0][0]
conv2d_7 (Conv2D)	(None, 32, 32, 64)	4160	max_pooling2d_1[0][0]
conv2d_8 (Conv2D)	(None, 32, 32, 192)	110784	conv2d_7[0][0]
max_pooling2d_2 (MaxPooling2D)	(None, 16, 16, 192)	0	conv2d_8[0][0]
conv2d_10 (Conv2D)	(None, 16, 16, 96)	18528	max_pooling2d_2[0][0]
conv2d_12 (Conv2D)	(None, 16, 16, 16)	3088	max_pooling2d_2[0][0]
max_pooling2d_3 (MaxPooling2D)	(None, 16, 16, 192)	0	max_pooling2d_2[0][0]
conv2d_9 (Conv2D)	(None, 16, 16, 64)	12352	max_pooling2d_2[0][0]
conv2d_11 (Conv2D)	(None, 16, 16, 128)	110720	conv2d_10[0][0]
conv2d_13 (Conv2D)	(None, 16, 16, 32)	12832	conv2d_12[0][0]
conv2d_14 (Conv2D)	(None, 16, 16, 32)	6176	

max_pooling2d_3[0][0]

concatenate_1 (Concatenate)	(None, 16, 16, 256)	0	conv2d_9[0][0] conv2d_11[0][0] conv2d_13[0][0] conv2d_14[0][0]
-----------------------------	---------------------	---	-------------------------------------------------------------------------

conv2d_16 (Conv2D)	(None, 16, 16, 128)	32896
concatenate_1[0][0]		

conv2d_18 (Conv2D)	(None, 16, 16, 32)	8224
concatenate_1[0][0]		

max_pooling2d_4 (MaxPooling2D)	(None, 16, 16, 256)	0
concatenate_1[0][0]		

conv2d_15 (Conv2D)	(None, 16, 16, 128)	32896
concatenate_1[0][0]		

conv2d_17 (Conv2D)	(None, 16, 16, 192)	221376	conv2d_16[0][0]
--------------------	---------------------	--------	-----------------

conv2d_19 (Conv2D)	(None, 16, 16, 96)	76896	conv2d_18[0][0]
--------------------	--------------------	-------	-----------------

conv2d_20 (Conv2D)	(None, 16, 16, 64)	16448
max_pooling2d_4[0][0]		

concatenate_2 (Concatenate)	(None, 16, 16, 480)	0	conv2d_15[0][0] conv2d_17[0][0] conv2d_19[0][0] conv2d_20[0][0]
-----------------------------	---------------------	---	--------------------------------------------------------------------------

max_pooling2d_5 (MaxPooling2D)	(None, 8, 8, 480)	0
concatenate_2[0][0]		

conv2d_22 (Conv2D)	(None, 8, 8, 96)	46176
max_pooling2d_5[0][0]		

conv2d_24 (Conv2D)	(None, 8, 8, 16)	7696	
max_pooling2d_5[0][0]			
max_pooling2d_6 (MaxPooling2D)	(None, 8, 8, 480)	0	
max_pooling2d_5[0][0]			
conv2d_21 (Conv2D)	(None, 8, 8, 192)	92352	
max_pooling2d_5[0][0]			
conv2d_23 (Conv2D)	(None, 8, 8, 208)	179920	conv2d_22[0][0]
conv2d_25 (Conv2D)	(None, 8, 8, 48)	19248	conv2d_24[0][0]
conv2d_26 (Conv2D)	(None, 8, 8, 64)	30784	
max_pooling2d_6[0][0]			
concatenate_3 (Concatenate)	(None, 8, 8, 512)	0	conv2d_21[0][0] conv2d_23[0][0] conv2d_25[0][0] conv2d_26[0][0]
conv2d_29 (Conv2D)	(None, 8, 8, 112)	57456	
concatenate_3[0][0]			
conv2d_31 (Conv2D)	(None, 8, 8, 24)	12312	
concatenate_3[0][0]			
max_pooling2d_7 (MaxPooling2D)	(None, 8, 8, 512)	0	
concatenate_3[0][0]			
conv2d_28 (Conv2D)	(None, 8, 8, 160)	82080	
concatenate_3[0][0]			
conv2d_30 (Conv2D)	(None, 8, 8, 224)	226016	conv2d_29[0][0]

conv2d_32 (Conv2D)	(None, 8, 8, 64)	38464	conv2d_31[0][0]

conv2d_33 (Conv2D)	(None, 8, 8, 64)	32832	
max_pooling2d_7[0][0]			

concatenate_4 (Concatenate)	(None, 8, 8, 512)	0	conv2d_28[0][0] conv2d_30[0][0] conv2d_32[0][0] conv2d_33[0][0]

conv2d_35 (Conv2D)	(None, 8, 8, 128)	65664	
concatenate_4[0][0]			

conv2d_37 (Conv2D)	(None, 8, 8, 24)	12312	
concatenate_4[0][0]			

max_pooling2d_8 (MaxPooling2D)	(None, 8, 8, 512)	0	
concatenate_4[0][0]			

conv2d_34 (Conv2D)	(None, 8, 8, 128)	65664	
concatenate_4[0][0]			

conv2d_36 (Conv2D)	(None, 8, 8, 256)	295168	conv2d_35[0][0]

conv2d_38 (Conv2D)	(None, 8, 8, 64)	38464	conv2d_37[0][0]

conv2d_39 (Conv2D)	(None, 8, 8, 64)	32832	
max_pooling2d_8[0][0]			

concatenate_5 (Concatenate)	(None, 8, 8, 512)	0	conv2d_34[0][0] conv2d_36[0][0] conv2d_38[0][0] conv2d_39[0][0]

conv2d_41 (Conv2D)	(None, 8, 8, 144)	73872	
concatenate_5[0][0]			

conv2d_43 (Conv2D)	(None, 8, 8, 32)	16416	
concatenate_5[0][0]			
max_pooling2d_9 (MaxPooling2D)	(None, 8, 8, 512)	0	
concatenate_5[0][0]			
conv2d_40 (Conv2D)	(None, 8, 8, 112)	57456	
concatenate_5[0][0]			
conv2d_42 (Conv2D)	(None, 8, 8, 288)	373536	conv2d_41[0][0]
conv2d_44 (Conv2D)	(None, 8, 8, 64)	51264	conv2d_43[0][0]
conv2d_45 (Conv2D)	(None, 8, 8, 64)	32832	
max_pooling2d_9[0][0]			
concatenate_6 (Concatenate)	(None, 8, 8, 528)	0	conv2d_40[0][0] conv2d_42[0][0] conv2d_44[0][0] conv2d_45[0][0]
conv2d_48 (Conv2D)	(None, 8, 8, 160)	84640	
concatenate_6[0][0]			
conv2d_50 (Conv2D)	(None, 8, 8, 32)	16928	
concatenate_6[0][0]			
max_pooling2d_10 (MaxPooling2D)	(None, 8, 8, 528)	0	
concatenate_6[0][0]			
conv2d_47 (Conv2D)	(None, 8, 8, 256)	135424	
concatenate_6[0][0]			
conv2d_49 (Conv2D)	(None, 8, 8, 320)	461120	conv2d_48[0][0]

conv2d_51 (Conv2D)	(None, 8, 8, 128)	102528	conv2d_50[0][0]

conv2d_52 (Conv2D)	(None, 8, 8, 128)	67712	
max_pooling2d_10[0][0]			

concatenate_7 (Concatenate)	(None, 8, 8, 832)	0	conv2d_47[0][0] conv2d_49[0][0] conv2d_51[0][0] conv2d_52[0][0]

max_pooling2d_11 (MaxPooling2D)	(None, 4, 4, 832)	0	
concatenate_7[0][0]			

conv2d_54 (Conv2D)	(None, 4, 4, 160)	133280	
max_pooling2d_11[0][0]			

conv2d_56 (Conv2D)	(None, 4, 4, 32)	26656	
max_pooling2d_11[0][0]			

max_pooling2d_12 (MaxPooling2D)	(None, 4, 4, 832)	0	
max_pooling2d_11[0][0]			

conv2d_53 (Conv2D)	(None, 4, 4, 256)	213248	
max_pooling2d_11[0][0]			

conv2d_55 (Conv2D)	(None, 4, 4, 320)	461120	conv2d_54[0][0]

conv2d_57 (Conv2D)	(None, 4, 4, 128)	102528	conv2d_56[0][0]

conv2d_58 (Conv2D)	(None, 4, 4, 128)	106624	
max_pooling2d_12[0][0]			

concatenate_8 (Concatenate)	(None, 4, 4, 832)	0	conv2d_53[0][0] conv2d_55[0][0] conv2d_57[0][0] conv2d_58[0][0]

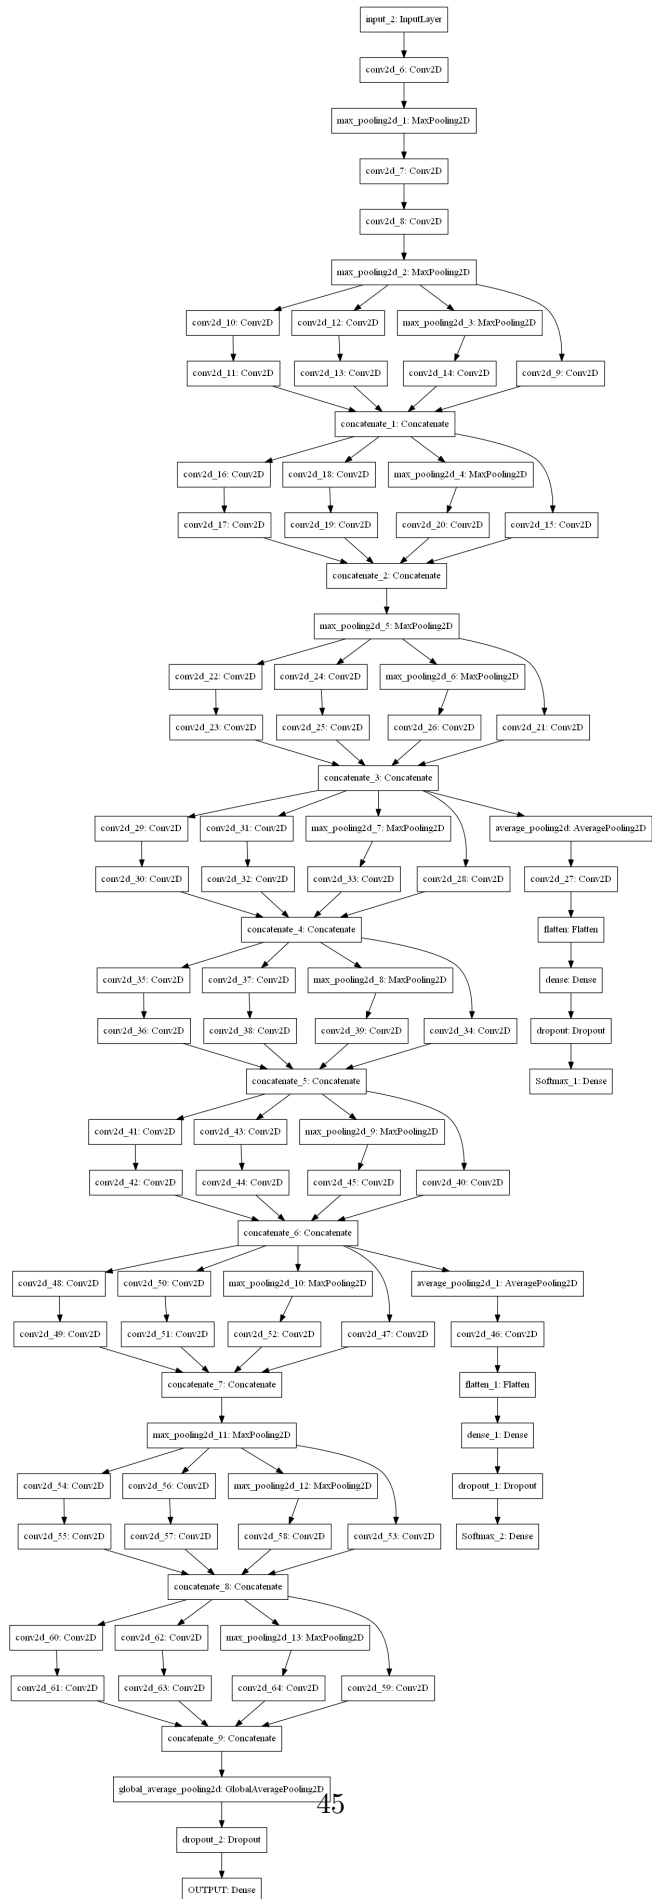
conv2d_60 (Conv2D)	(None, 4, 4, 192)	159936	
concatenate_8[0][0]			
conv2d_62 (Conv2D)	(None, 4, 4, 48)	39984	
concatenate_8[0][0]			
max_pooling2d_13 (MaxPooling2D)	(None, 4, 4, 832)	0	
concatenate_8[0][0]			
average_pooling2d (AveragePooli	(None, 2, 2, 512)	0	
concatenate_3[0][0]			
average_pooling2d_1 (AveragePoo	(None, 2, 2, 528)	0	
concatenate_6[0][0]			
conv2d_59 (Conv2D)	(None, 4, 4, 384)	319872	
concatenate_8[0][0]			
conv2d_61 (Conv2D)	(None, 4, 4, 384)	663936	conv2d_60[0][0]
conv2d_63 (Conv2D)	(None, 4, 4, 128)	153728	conv2d_62[0][0]
conv2d_64 (Conv2D)	(None, 4, 4, 128)	106624	
max_pooling2d_13[0][0]			
conv2d_27 (Conv2D)	(None, 2, 2, 128)	65664	
average_pooling2d[0][0]			
conv2d_46 (Conv2D)	(None, 2, 2, 128)	67712	
average_pooling2d_1[0][0]			
concatenate_9 (Concatenate)	(None, 4, 4, 1024)	0	conv2d_59[0][0] conv2d_61[0][0] conv2d_63[0][0] conv2d_64[0][0]

```

-----
flatten (Flatten)                (None, 512)          0          conv2d_27[0] [0]
-----
flatten_1 (Flatten)              (None, 512)          0          conv2d_46[0] [0]
-----
global_average_pooling2d (Globa (None, 1024)        0
concatenate_9[0] [0]
-----
dense (Dense)                    (None, 1024)        525312     flatten[0] [0]
-----
dense_1 (Dense)                  (None, 1024)        525312     flatten_1[0] [0]
-----
dropout_2 (Dropout)              (None, 1024)        0
global_average_pooling2d[0] [0]
-----
dropout (Dropout)                (None, 1024)        0          dense[0] [0]
-----
dropout_1 (Dropout)              (None, 1024)        0          dense_1[0] [0]
-----
OUTPUT (Dense)                  (None, 8)           8200       dropout_2[0] [0]
-----
Softmax_1 (Dense)                (None, 8)           8200       dropout[0] [0]
-----
Softmax_2 (Dense)                (None, 8)           8200       dropout_1[0] [0]
=====
Total params: 7,182,152
Trainable params: 7,182,152
Non-trainable params: 0
-----

```

[12]:



```
[13]: epochs = 50
initial_lr = 0.01

def decay(epoch, steps=100):
    initial_lr = 0.01
    drop = 0.96
    epochs_drop = 8
    lr = initial_lr * math.pow(drop, math.floor((1+epoch)/epochs_drop))
    return lr

sgd = SGD(lr=initial_lr, momentum=0.9, nesterov=False)

lr_scheduler = LearningRateScheduler(decay, verbose=1)

model.compile(loss=['categorical_crossentropy', 'categorical_crossentropy', 'categorical_crossentropy'],
              optimizer=sgd, metrics=['accuracy'])

[14]: history = model.fit(train_images, [train_labels, train_labels, train_labels],
                        validation_data=(test_images, [test_labels, test_labels, test_labels]),
                        epochs=epochs, batch_size=256, callbacks=[lr_scheduler])
```

```
Epoch 00001: LearningRateScheduler reducing learning rate to 0.01.
Epoch 1/50
 2/40 [>...] - ETA: 28s - loss: 6.2378 - OUTPUT_loss:
2.0805 - Softmax_1_loss: 2.0780 - Softmax_2_loss: 2.0793 - OUTPUT_accuracy:
0.1113 - Softmax_1_accuracy: 0.1309 - Softmax_2_accuracy:
0.1016WARNING:tensorflow:Callbacks method `on_train_batch_end` is slow compared
to the batch time (batch time: 0.5235s vs `on_train_batch_end` time: 0.9779s).
Check your callbacks.
40/40 [=====] - 76s 2s/step - loss: 6.2381 -
OUTPUT_loss: 2.0800 - Softmax_1_loss: 2.0787 - Softmax_2_loss: 2.0795 -
OUTPUT_accuracy: 0.1202 - Softmax_1_accuracy: 0.1354 - Softmax_2_accuracy:
0.1252 - val_loss: 6.2335 - val_OUTPUT_loss: 2.0794 - val_Softmax_1_loss: 2.0748
- val_Softmax_2_loss: 2.0793 - val_OUTPUT_accuracy: 0.1195 -
val_Softmax_1_accuracy: 0.1925 - val_Softmax_2_accuracy: 0.1195

Epoch 00002: LearningRateScheduler reducing learning rate to 0.01.
Epoch 2/50
40/40 [=====] - 58s 1s/step - loss: 6.2260 -
OUTPUT_loss: 2.0785 - Softmax_1_loss: 2.0700 - Softmax_2_loss: 2.0775 -
OUTPUT_accuracy: 0.1307 - Softmax_1_accuracy: 0.1569 - Softmax_2_accuracy:
0.1378 - val_loss: 6.2009 - val_OUTPUT_loss: 2.0758 - val_Softmax_1_loss: 2.0516
- val_Softmax_2_loss: 2.0734 - val_OUTPUT_accuracy: 0.2115 -
```

val_Softmax_1_accuracy: 0.2090 - val_Softmax_2_accuracy: 0.1990

Epoch 00003: LearningRateScheduler reducing learning rate to 0.01.

Epoch 3/50

40/40 [=====] - 60s 2s/step - loss: 6.1570 -
OUTPUT_loss: 2.0654 - Softmax_1_loss: 2.0314 - Softmax_2_loss: 2.0602 -
OUTPUT_accuracy: 0.1744 - Softmax_1_accuracy: 0.1802 - Softmax_2_accuracy:
0.1612 - val_loss: 6.0465 - val_OUTPUT_loss: 2.0305 - val_Softmax_1_loss: 2.0004
- val_Softmax_2_loss: 2.0156 - val_OUTPUT_accuracy: 0.2305 -
val_Softmax_1_accuracy: 0.2185 - val_Softmax_2_accuracy: 0.2175

Epoch 00004: LearningRateScheduler reducing learning rate to 0.01.

Epoch 4/50

40/40 [=====] - 61s 2s/step - loss: 5.8924 -
OUTPUT_loss: 1.9629 - Softmax_1_loss: 1.9611 - Softmax_2_loss: 1.9683 -
OUTPUT_accuracy: 0.2097 - Softmax_1_accuracy: 0.2078 - Softmax_2_accuracy:
0.2101 - val_loss: 5.7264 - val_OUTPUT_loss: 1.9135 - val_Softmax_1_loss: 1.9060
- val_Softmax_2_loss: 1.9068 - val_OUTPUT_accuracy: 0.2400 -
val_Softmax_1_accuracy: 0.2655 - val_Softmax_2_accuracy: 0.2410

Epoch 00005: LearningRateScheduler reducing learning rate to 0.01.

Epoch 5/50

40/40 [=====] - 59s 1s/step - loss: 5.7030 -
OUTPUT_loss: 1.8967 - Softmax_1_loss: 1.8983 - Softmax_2_loss: 1.9080 -
OUTPUT_accuracy: 0.2458 - Softmax_1_accuracy: 0.2504 - Softmax_2_accuracy:
0.2338 - val_loss: 5.8739 - val_OUTPUT_loss: 1.9272 - val_Softmax_1_loss: 2.0244
- val_Softmax_2_loss: 1.9224 - val_OUTPUT_accuracy: 0.2025 -
val_Softmax_1_accuracy: 0.1935 - val_Softmax_2_accuracy: 0.1980

Epoch 00006: LearningRateScheduler reducing learning rate to 0.01.

Epoch 6/50

40/40 [=====] - 60s 2s/step - loss: 5.6529 -
OUTPUT_loss: 1.8677 - Softmax_1_loss: 1.9008 - Softmax_2_loss: 1.8844 -
OUTPUT_accuracy: 0.2480 - Softmax_1_accuracy: 0.2390 - Softmax_2_accuracy:
0.2472 - val_loss: 5.3931 - val_OUTPUT_loss: 1.7901 - val_Softmax_1_loss: 1.8033
- val_Softmax_2_loss: 1.7997 - val_OUTPUT_accuracy: 0.2950 -
val_Softmax_1_accuracy: 0.3110 - val_Softmax_2_accuracy: 0.2820

Epoch 00007: LearningRateScheduler reducing learning rate to 0.01.

Epoch 7/50

40/40 [=====] - 32s 807ms/step - loss: 5.4414 -
OUTPUT_loss: 1.8105 - Softmax_1_loss: 1.8012 - Softmax_2_loss: 1.8297 -
OUTPUT_accuracy: 0.2816 - Softmax_1_accuracy: 0.2970 - Softmax_2_accuracy:
0.2775 - val_loss: 5.6986 - val_OUTPUT_loss: 1.9106 - val_Softmax_1_loss: 1.8835
- val_Softmax_2_loss: 1.9044 - val_OUTPUT_accuracy: 0.2430 -
val_Softmax_1_accuracy: 0.2460 - val_Softmax_2_accuracy: 0.2455

Epoch 00008: LearningRateScheduler reducing learning rate to 0.0096.

Epoch 8/50
40/40 [=====] - 22s 560ms/step - loss: 5.4563 -
OUTPUT_loss: 1.8184 - Softmax_1_loss: 1.8176 - Softmax_2_loss: 1.8203 -
OUTPUT_accuracy: 0.2891 - Softmax_1_accuracy: 0.2973 - Softmax_2_accuracy:
0.2905 - val_loss: 5.2625 - val_OUTPUT_loss: 1.7581 - val_Softmax_1_loss: 1.7520
- val_Softmax_2_loss: 1.7525 - val_OUTPUT_accuracy: 0.3075 -
val_Softmax_1_accuracy: 0.3220 - val_Softmax_2_accuracy: 0.3195

Epoch 00009: LearningRateScheduler reducing learning rate to 0.0096.

Epoch 9/50
40/40 [=====] - 23s 563ms/step - loss: 5.0586 -
OUTPUT_loss: 1.6841 - Softmax_1_loss: 1.6824 - Softmax_2_loss: 1.6921 -
OUTPUT_accuracy: 0.3366 - Softmax_1_accuracy: 0.3406 - Softmax_2_accuracy:
0.3318 - val_loss: 5.6729 - val_OUTPUT_loss: 1.9650 - val_Softmax_1_loss: 1.7989
- val_Softmax_2_loss: 1.9090 - val_OUTPUT_accuracy: 0.2625 -
val_Softmax_1_accuracy: 0.3205 - val_Softmax_2_accuracy: 0.2845

Epoch 00010: LearningRateScheduler reducing learning rate to 0.0096.

Epoch 10/50
40/40 [=====] - 23s 563ms/step - loss: 4.9203 -
OUTPUT_loss: 1.6403 - Softmax_1_loss: 1.6307 - Softmax_2_loss: 1.6493 -
OUTPUT_accuracy: 0.3594 - Softmax_1_accuracy: 0.3631 - Softmax_2_accuracy:
0.3564 - val_loss: 5.1822 - val_OUTPUT_loss: 1.7407 - val_Softmax_1_loss: 1.7411
- val_Softmax_2_loss: 1.7004 - val_OUTPUT_accuracy: 0.3240 -
val_Softmax_1_accuracy: 0.3525 - val_Softmax_2_accuracy: 0.3425

Epoch 00011: LearningRateScheduler reducing learning rate to 0.0096.

Epoch 11/50
40/40 [=====] - 23s 565ms/step - loss: 4.7377 -
OUTPUT_loss: 1.5797 - Softmax_1_loss: 1.5774 - Softmax_2_loss: 1.5805 -
OUTPUT_accuracy: 0.3862 - Softmax_1_accuracy: 0.3932 - Softmax_2_accuracy:
0.3872 - val_loss: 5.1441 - val_OUTPUT_loss: 1.7678 - val_Softmax_1_loss: 1.6014
- val_Softmax_2_loss: 1.7749 - val_OUTPUT_accuracy: 0.3085 -
val_Softmax_1_accuracy: 0.3730 - val_Softmax_2_accuracy: 0.3130

Epoch 00012: LearningRateScheduler reducing learning rate to 0.0096.

Epoch 12/50
40/40 [=====] - 23s 564ms/step - loss: 4.4626 -
OUTPUT_loss: 1.4812 - Softmax_1_loss: 1.4931 - Softmax_2_loss: 1.4882 -
OUTPUT_accuracy: 0.4222 - Softmax_1_accuracy: 0.4213 - Softmax_2_accuracy:
0.4202 - val_loss: 5.8929 - val_OUTPUT_loss: 2.0213 - val_Softmax_1_loss: 1.8636
- val_Softmax_2_loss: 2.0080 - val_OUTPUT_accuracy: 0.2185 -
val_Softmax_1_accuracy: 0.2680 - val_Softmax_2_accuracy: 0.2275

Epoch 00013: LearningRateScheduler reducing learning rate to 0.0096.

Epoch 13/50
40/40 [=====] - 23s 564ms/step - loss: 5.2081 -
OUTPUT_loss: 1.7512 - Softmax_1_loss: 1.7089 - Softmax_2_loss: 1.7479 -

OUTPUT_accuracy: 0.3144 - Softmax_1_accuracy: 0.3289 - Softmax_2_accuracy:
0.3141 - val_loss: 5.0090 - val_OUTPUT_loss: 1.7137 - val_Softmax_1_loss: 1.6147
- val_Softmax_2_loss: 1.6806 - val_OUTPUT_accuracy: 0.3580 -
val_Softmax_1_accuracy: 0.3925 - val_Softmax_2_accuracy: 0.3740

Epoch 00014: LearningRateScheduler reducing learning rate to 0.0096.

Epoch 14/50

40/40 [=====] - 23s 566ms/step - loss: 4.4792 -
OUTPUT_loss: 1.4936 - Softmax_1_loss: 1.4854 - Softmax_2_loss: 1.5002 -
OUTPUT_accuracy: 0.4258 - Softmax_1_accuracy: 0.4297 - Softmax_2_accuracy:
0.4236 - val_loss: 7.5746 - val_OUTPUT_loss: 2.5663 - val_Softmax_1_loss: 2.6947
- val_Softmax_2_loss: 2.3136 - val_OUTPUT_accuracy: 0.1910 -
val_Softmax_1_accuracy: 0.1925 - val_Softmax_2_accuracy: 0.2095

Epoch 00015: LearningRateScheduler reducing learning rate to 0.0096.

Epoch 15/50

40/40 [=====] - 23s 564ms/step - loss: 4.9839 -
OUTPUT_loss: 1.6537 - Softmax_1_loss: 1.6829 - Softmax_2_loss: 1.6473 -
OUTPUT_accuracy: 0.3685 - Softmax_1_accuracy: 0.3582 - Softmax_2_accuracy:
0.3707 - val_loss: 4.5518 - val_OUTPUT_loss: 1.5021 - val_Softmax_1_loss: 1.5368
- val_Softmax_2_loss: 1.5129 - val_OUTPUT_accuracy: 0.4095 -
val_Softmax_1_accuracy: 0.4160 - val_Softmax_2_accuracy: 0.4090

Epoch 00016: LearningRateScheduler reducing learning rate to 0.009216.

Epoch 16/50

40/40 [=====] - 23s 566ms/step - loss: 4.3151 -
OUTPUT_loss: 1.4291 - Softmax_1_loss: 1.4450 - Softmax_2_loss: 1.4410 -
OUTPUT_accuracy: 0.4376 - Softmax_1_accuracy: 0.4355 - Softmax_2_accuracy:
0.4338 - val_loss: 5.2748 - val_OUTPUT_loss: 1.8029 - val_Softmax_1_loss: 1.7123
- val_Softmax_2_loss: 1.7596 - val_OUTPUT_accuracy: 0.3380 -
val_Softmax_1_accuracy: 0.3580 - val_Softmax_2_accuracy: 0.3545

Epoch 00017: LearningRateScheduler reducing learning rate to 0.009216.

Epoch 17/50

40/40 [=====] - 23s 566ms/step - loss: 4.2401 -
OUTPUT_loss: 1.4129 - Softmax_1_loss: 1.4116 - Softmax_2_loss: 1.4155 -
OUTPUT_accuracy: 0.4531 - Softmax_1_accuracy: 0.4570 - Softmax_2_accuracy:
0.4531 - val_loss: 4.2604 - val_OUTPUT_loss: 1.3928 - val_Softmax_1_loss: 1.4321
- val_Softmax_2_loss: 1.4355 - val_OUTPUT_accuracy: 0.4565 -
val_Softmax_1_accuracy: 0.4630 - val_Softmax_2_accuracy: 0.4545

Epoch 00018: LearningRateScheduler reducing learning rate to 0.009216.

Epoch 18/50

40/40 [=====] - 23s 565ms/step - loss: 3.9672 -
OUTPUT_loss: 1.3237 - Softmax_1_loss: 1.3192 - Softmax_2_loss: 1.3244 -
OUTPUT_accuracy: 0.4833 - Softmax_1_accuracy: 0.4916 - Softmax_2_accuracy:
0.4825 - val_loss: 4.3107 - val_OUTPUT_loss: 1.4450 - val_Softmax_1_loss: 1.4249
- val_Softmax_2_loss: 1.4409 - val_OUTPUT_accuracy: 0.4270 -

val_Softmax_1_accuracy: 0.4245 - val_Softmax_2_accuracy: 0.4275

Epoch 00019: LearningRateScheduler reducing learning rate to 0.009216.

Epoch 19/50

40/40 [=====] - 23s 565ms/step - loss: 3.8298 -
OUTPUT_loss: 1.2729 - Softmax_1_loss: 1.2791 - Softmax_2_loss: 1.2778 -
OUTPUT_accuracy: 0.5065 - Softmax_1_accuracy: 0.5143 - Softmax_2_accuracy:
0.5077 - val_loss: 4.2529 - val_OUTPUT_loss: 1.4734 - val_Softmax_1_loss: 1.3554
- val_Softmax_2_loss: 1.4242 - val_OUTPUT_accuracy: 0.4435 -
val_Softmax_1_accuracy: 0.4835 - val_Softmax_2_accuracy: 0.4705

Epoch 00020: LearningRateScheduler reducing learning rate to 0.009216.

Epoch 20/50

40/40 [=====] - 23s 568ms/step - loss: 3.8400 -
OUTPUT_loss: 1.2779 - Softmax_1_loss: 1.2700 - Softmax_2_loss: 1.2920 -
OUTPUT_accuracy: 0.5160 - Softmax_1_accuracy: 0.5203 - Softmax_2_accuracy:
0.5107 - val_loss: 3.8997 - val_OUTPUT_loss: 1.3221 - val_Softmax_1_loss: 1.2833
- val_Softmax_2_loss: 1.2943 - val_OUTPUT_accuracy: 0.5000 -
val_Softmax_1_accuracy: 0.5130 - val_Softmax_2_accuracy: 0.5100

Epoch 00021: LearningRateScheduler reducing learning rate to 0.009216.

Epoch 21/50

40/40 [=====] - 23s 569ms/step - loss: 3.6066 -
OUTPUT_loss: 1.2023 - Softmax_1_loss: 1.1995 - Softmax_2_loss: 1.2049 -
OUTPUT_accuracy: 0.5367 - Softmax_1_accuracy: 0.5434 - Softmax_2_accuracy:
0.5371 - val_loss: 3.7307 - val_OUTPUT_loss: 1.2403 - val_Softmax_1_loss: 1.2428
- val_Softmax_2_loss: 1.2476 - val_OUTPUT_accuracy: 0.5355 -
val_Softmax_1_accuracy: 0.5375 - val_Softmax_2_accuracy: 0.5410

Epoch 00022: LearningRateScheduler reducing learning rate to 0.009216.

Epoch 22/50

40/40 [=====] - 23s 567ms/step - loss: 3.4835 -
OUTPUT_loss: 1.1582 - Softmax_1_loss: 1.1588 - Softmax_2_loss: 1.1665 -
OUTPUT_accuracy: 0.5573 - Softmax_1_accuracy: 0.5576 - Softmax_2_accuracy:
0.5543 - val_loss: 4.0985 - val_OUTPUT_loss: 1.4268 - val_Softmax_1_loss: 1.2844
- val_Softmax_2_loss: 1.3874 - val_OUTPUT_accuracy: 0.4570 -
val_Softmax_1_accuracy: 0.4930 - val_Softmax_2_accuracy: 0.4695

Epoch 00023: LearningRateScheduler reducing learning rate to 0.009216.

Epoch 23/50

40/40 [=====] - 23s 568ms/step - loss: 3.4658 -
OUTPUT_loss: 1.1523 - Softmax_1_loss: 1.1570 - Softmax_2_loss: 1.1565 -
OUTPUT_accuracy: 0.5618 - Softmax_1_accuracy: 0.5581 - Softmax_2_accuracy:
0.5597 - val_loss: 3.4725 - val_OUTPUT_loss: 1.1674 - val_Softmax_1_loss: 1.1540
- val_Softmax_2_loss: 1.1511 - val_OUTPUT_accuracy: 0.5660 -
val_Softmax_1_accuracy: 0.5785 - val_Softmax_2_accuracy: 0.5825

Epoch 00024: LearningRateScheduler reducing learning rate to

0.008847359999999999.

Epoch 24/50

40/40 [=====] - 23s 566ms/step - loss: 3.2603 -
OUTPUT_loss: 1.0851 - Softmax_1_loss: 1.0856 - Softmax_2_loss: 1.0897 -
OUTPUT_accuracy: 0.5924 - Softmax_1_accuracy: 0.5889 - Softmax_2_accuracy:
0.5876 - val_loss: 4.4006 - val_OUTPUT_loss: 1.5657 - val_Softmax_1_loss: 1.3881
- val_Softmax_2_loss: 1.4468 - val_OUTPUT_accuracy: 0.4600 -
val_Softmax_1_accuracy: 0.4945 - val_Softmax_2_accuracy: 0.4820

Epoch 00025: LearningRateScheduler reducing learning rate to

0.008847359999999999.

Epoch 25/50

40/40 [=====] - 23s 567ms/step - loss: 3.9236 -
OUTPUT_loss: 1.3001 - Softmax_1_loss: 1.3077 - Softmax_2_loss: 1.3157 -
OUTPUT_accuracy: 0.5050 - Softmax_1_accuracy: 0.5059 - Softmax_2_accuracy:
0.5027 - val_loss: 3.4034 - val_OUTPUT_loss: 1.1397 - val_Softmax_1_loss: 1.1338
- val_Softmax_2_loss: 1.1299 - val_OUTPUT_accuracy: 0.5745 -
val_Softmax_1_accuracy: 0.5790 - val_Softmax_2_accuracy: 0.5715

Epoch 00026: LearningRateScheduler reducing learning rate to

0.008847359999999999.

Epoch 26/50

40/40 [=====] - 23s 567ms/step - loss: 3.2459 -
OUTPUT_loss: 1.0798 - Softmax_1_loss: 1.0802 - Softmax_2_loss: 1.0859 -
OUTPUT_accuracy: 0.5888 - Softmax_1_accuracy: 0.5891 - Softmax_2_accuracy:
0.5856 - val_loss: 3.4042 - val_OUTPUT_loss: 1.1540 - val_Softmax_1_loss: 1.1075
- val_Softmax_2_loss: 1.1426 - val_OUTPUT_accuracy: 0.5575 -
val_Softmax_1_accuracy: 0.5870 - val_Softmax_2_accuracy: 0.5645

Epoch 00027: LearningRateScheduler reducing learning rate to

0.008847359999999999.

Epoch 27/50

40/40 [=====] - 23s 568ms/step - loss: 3.0722 -
OUTPUT_loss: 1.0155 - Softmax_1_loss: 1.0257 - Softmax_2_loss: 1.0310 -
OUTPUT_accuracy: 0.6114 - Softmax_1_accuracy: 0.6141 - Softmax_2_accuracy:
0.6105 - val_loss: 3.2974 - val_OUTPUT_loss: 1.0850 - val_Softmax_1_loss: 1.1174
- val_Softmax_2_loss: 1.0949 - val_OUTPUT_accuracy: 0.5865 -
val_Softmax_1_accuracy: 0.5755 - val_Softmax_2_accuracy: 0.5890

Epoch 00028: LearningRateScheduler reducing learning rate to

0.008847359999999999.

Epoch 28/50

40/40 [=====] - 23s 569ms/step - loss: 2.9450 -
OUTPUT_loss: 0.9732 - Softmax_1_loss: 0.9897 - Softmax_2_loss: 0.9821 -
OUTPUT_accuracy: 0.6333 - Softmax_1_accuracy: 0.6281 - Softmax_2_accuracy:
0.6334 - val_loss: 4.0322 - val_OUTPUT_loss: 1.3557 - val_Softmax_1_loss: 1.3315
- val_Softmax_2_loss: 1.3450 - val_OUTPUT_accuracy: 0.5175 -
val_Softmax_1_accuracy: 0.5300 - val_Softmax_2_accuracy: 0.5315

Epoch 00029: LearningRateScheduler reducing learning rate to 0.008847359999999999.

Epoch 29/50

40/40 [=====] - 23s 567ms/step - loss: 3.0890 -
OUTPUT_loss: 1.0268 - Softmax_1_loss: 1.0290 - Softmax_2_loss: 1.0333 -
OUTPUT_accuracy: 0.6241 - Softmax_1_accuracy: 0.6233 - Softmax_2_accuracy:
0.6194 - val_loss: 3.7559 - val_OUTPUT_loss: 1.2649 - val_Softmax_1_loss: 1.2274
- val_Softmax_2_loss: 1.2637 - val_OUTPUT_accuracy: 0.5310 -
val_Softmax_1_accuracy: 0.5525 - val_Softmax_2_accuracy: 0.5330

Epoch 00030: LearningRateScheduler reducing learning rate to 0.008847359999999999.

Epoch 30/50

40/40 [=====] - 23s 567ms/step - loss: 2.8281 -
OUTPUT_loss: 0.9301 - Softmax_1_loss: 0.9531 - Softmax_2_loss: 0.9450 -
OUTPUT_accuracy: 0.6531 - Softmax_1_accuracy: 0.6467 - Softmax_2_accuracy:
0.6504 - val_loss: 2.8133 - val_OUTPUT_loss: 0.9415 - val_Softmax_1_loss: 0.9289
- val_Softmax_2_loss: 0.9428 - val_OUTPUT_accuracy: 0.6470 -
val_Softmax_1_accuracy: 0.6615 - val_Softmax_2_accuracy: 0.6460

Epoch 00031: LearningRateScheduler reducing learning rate to 0.008847359999999999.

Epoch 31/50

40/40 [=====] - 23s 567ms/step - loss: 2.7152 -
OUTPUT_loss: 0.8959 - Softmax_1_loss: 0.9133 - Softmax_2_loss: 0.9060 -
OUTPUT_accuracy: 0.6670 - Softmax_1_accuracy: 0.6640 - Softmax_2_accuracy:
0.6608 - val_loss: 3.5656 - val_OUTPUT_loss: 1.1958 - val_Softmax_1_loss: 1.1756
- val_Softmax_2_loss: 1.1943 - val_OUTPUT_accuracy: 0.5580 -
val_Softmax_1_accuracy: 0.5680 - val_Softmax_2_accuracy: 0.5580

Epoch 00032: LearningRateScheduler reducing learning rate to 0.008493465599999998.

Epoch 32/50

40/40 [=====] - 23s 568ms/step - loss: 2.7815 -
OUTPUT_loss: 0.9190 - Softmax_1_loss: 0.9298 - Softmax_2_loss: 0.9328 -
OUTPUT_accuracy: 0.6575 - Softmax_1_accuracy: 0.6515 - Softmax_2_accuracy:
0.6549 - val_loss: 3.3008 - val_OUTPUT_loss: 1.1112 - val_Softmax_1_loss: 1.1035
- val_Softmax_2_loss: 1.0860 - val_OUTPUT_accuracy: 0.6080 -
val_Softmax_1_accuracy: 0.6020 - val_Softmax_2_accuracy: 0.6140

Epoch 00033: LearningRateScheduler reducing learning rate to 0.008493465599999998.

Epoch 33/50

40/40 [=====] - 23s 568ms/step - loss: 2.6493 -
OUTPUT_loss: 0.8755 - Softmax_1_loss: 0.8880 - Softmax_2_loss: 0.8858 -
OUTPUT_accuracy: 0.6739 - Softmax_1_accuracy: 0.6692 - Softmax_2_accuracy:
0.6676 - val_loss: 6.0949 - val_OUTPUT_loss: 2.1007 - val_Softmax_1_loss: 1.8723

- val_Softmax_2_loss: 2.1218 - val_OUTPUT_accuracy: 0.3770 -
val_Softmax_1_accuracy: 0.4125 - val_Softmax_2_accuracy: 0.3840

Epoch 00034: LearningRateScheduler reducing learning rate to
0.008493465599999998.

Epoch 34/50

40/40 [=====] - 23s 566ms/step - loss: 3.4261 -
OUTPUT_loss: 1.1471 - Softmax_1_loss: 1.1279 - Softmax_2_loss: 1.1511 -
OUTPUT_accuracy: 0.5782 - Softmax_1_accuracy: 0.5892 - Softmax_2_accuracy:
0.5754 - val_loss: 3.4782 - val_OUTPUT_loss: 1.1641 - val_Softmax_1_loss: 1.1742
- val_Softmax_2_loss: 1.1398 - val_OUTPUT_accuracy: 0.5715 -
val_Softmax_1_accuracy: 0.5760 - val_Softmax_2_accuracy: 0.5820

Epoch 00035: LearningRateScheduler reducing learning rate to
0.008493465599999998.

Epoch 35/50

40/40 [=====] - 23s 567ms/step - loss: 2.6128 -
OUTPUT_loss: 0.8586 - Softmax_1_loss: 0.8836 - Softmax_2_loss: 0.8707 -
OUTPUT_accuracy: 0.6829 - Softmax_1_accuracy: 0.6766 - Softmax_2_accuracy:
0.6794 - val_loss: 3.2507 - val_OUTPUT_loss: 1.0843 - val_Softmax_1_loss: 1.0747
- val_Softmax_2_loss: 1.0918 - val_OUTPUT_accuracy: 0.6065 -
val_Softmax_1_accuracy: 0.6170 - val_Softmax_2_accuracy: 0.5990

Epoch 00036: LearningRateScheduler reducing learning rate to
0.008493465599999998.

Epoch 36/50

40/40 [=====] - 23s 568ms/step - loss: 2.5633 -
OUTPUT_loss: 0.8426 - Softmax_1_loss: 0.8609 - Softmax_2_loss: 0.8598 -
OUTPUT_accuracy: 0.6843 - Softmax_1_accuracy: 0.6792 - Softmax_2_accuracy:
0.6829 - val_loss: 2.6859 - val_OUTPUT_loss: 0.8918 - val_Softmax_1_loss: 0.9061
- val_Softmax_2_loss: 0.8880 - val_OUTPUT_accuracy: 0.6695 -
val_Softmax_1_accuracy: 0.6775 - val_Softmax_2_accuracy: 0.6720

Epoch 00037: LearningRateScheduler reducing learning rate to
0.008493465599999998.

Epoch 37/50

40/40 [=====] - 23s 567ms/step - loss: 2.4281 -
OUTPUT_loss: 0.7974 - Softmax_1_loss: 0.8189 - Softmax_2_loss: 0.8117 -
OUTPUT_accuracy: 0.7030 - Softmax_1_accuracy: 0.6980 - Softmax_2_accuracy:
0.6998 - val_loss: 3.6172 - val_OUTPUT_loss: 1.2336 - val_Softmax_1_loss: 1.1798
- val_Softmax_2_loss: 1.2038 - val_OUTPUT_accuracy: 0.5990 -
val_Softmax_1_accuracy: 0.5940 - val_Softmax_2_accuracy: 0.5995

Epoch 00038: LearningRateScheduler reducing learning rate to
0.008493465599999998.

Epoch 38/50

40/40 [=====] - 23s 567ms/step - loss: 2.4884 -
OUTPUT_loss: 0.8209 - Softmax_1_loss: 0.8345 - Softmax_2_loss: 0.8330 -

OUTPUT_accuracy: 0.7002 - Softmax_1_accuracy: 0.6952 - Softmax_2_accuracy:
0.6924 - val_loss: 3.1439 - val_OUTPUT_loss: 1.0982 - val_Softmax_1_loss: 0.9850
- val_Softmax_2_loss: 1.0608 - val_OUTPUT_accuracy: 0.5930 -
val_Softmax_1_accuracy: 0.6320 - val_Softmax_2_accuracy: 0.6000

Epoch 00039: LearningRateScheduler reducing learning rate to
0.008493465599999998.

Epoch 39/50

40/40 [=====] - 23s 568ms/step - loss: 2.3307 -
OUTPUT_loss: 0.7709 - Softmax_1_loss: 0.7849 - Softmax_2_loss: 0.7749 -
OUTPUT_accuracy: 0.7084 - Softmax_1_accuracy: 0.7051 - Softmax_2_accuracy:
0.7086 - val_loss: 2.9164 - val_OUTPUT_loss: 0.9848 - val_Softmax_1_loss: 0.9352
- val_Softmax_2_loss: 0.9964 - val_OUTPUT_accuracy: 0.6435 -
val_Softmax_1_accuracy: 0.6600 - val_Softmax_2_accuracy: 0.6375

Epoch 00040: LearningRateScheduler reducing learning rate to 0.008153726976.

Epoch 40/50

40/40 [=====] - 23s 572ms/step - loss: 2.2766 -
OUTPUT_loss: 0.7438 - Softmax_1_loss: 0.7699 - Softmax_2_loss: 0.7629 -
OUTPUT_accuracy: 0.7220 - Softmax_1_accuracy: 0.7161 - Softmax_2_accuracy:
0.7177 - val_loss: 2.4460 - val_OUTPUT_loss: 0.8220 - val_Softmax_1_loss: 0.8140
- val_Softmax_2_loss: 0.8101 - val_OUTPUT_accuracy: 0.6985 -
val_Softmax_1_accuracy: 0.6980 - val_Softmax_2_accuracy: 0.6995

Epoch 00041: LearningRateScheduler reducing learning rate to 0.008153726976.

Epoch 41/50

40/40 [=====] - 23s 574ms/step - loss: 2.0873 -
OUTPUT_loss: 0.6809 - Softmax_1_loss: 0.7163 - Softmax_2_loss: 0.6901 -
OUTPUT_accuracy: 0.7494 - Softmax_1_accuracy: 0.7368 - Softmax_2_accuracy:
0.7433 - val_loss: 2.8081 - val_OUTPUT_loss: 0.9376 - val_Softmax_1_loss: 0.9417
- val_Softmax_2_loss: 0.9288 - val_OUTPUT_accuracy: 0.6610 -
val_Softmax_1_accuracy: 0.6730 - val_Softmax_2_accuracy: 0.6640

Epoch 00042: LearningRateScheduler reducing learning rate to 0.008153726976.

Epoch 42/50

40/40 [=====] - 23s 572ms/step - loss: 2.2156 -
OUTPUT_loss: 0.7238 - Softmax_1_loss: 0.7553 - Softmax_2_loss: 0.7365 -
OUTPUT_accuracy: 0.7334 - Softmax_1_accuracy: 0.7235 - Softmax_2_accuracy:
0.7288 - val_loss: 2.5242 - val_OUTPUT_loss: 0.8492 - val_Softmax_1_loss: 0.8333
- val_Softmax_2_loss: 0.8417 - val_OUTPUT_accuracy: 0.6875 -
val_Softmax_1_accuracy: 0.7015 - val_Softmax_2_accuracy: 0.6890

Epoch 00043: LearningRateScheduler reducing learning rate to 0.008153726976.

Epoch 43/50

40/40 [=====] - 23s 567ms/step - loss: 1.9847 -
OUTPUT_loss: 0.6456 - Softmax_1_loss: 0.6779 - Softmax_2_loss: 0.6611 -
OUTPUT_accuracy: 0.7564 - Softmax_1_accuracy: 0.7515 - Softmax_2_accuracy:
0.7564 - val_loss: 2.5666 - val_OUTPUT_loss: 0.8688 - val_Softmax_1_loss: 0.8352

- val_Softmax_2_loss: 0.8626 - val_OUTPUT_accuracy: 0.6985 -
val_Softmax_1_accuracy: 0.6950 - val_Softmax_2_accuracy: 0.7000

Epoch 00044: LearningRateScheduler reducing learning rate to 0.008153726976.

Epoch 44/50

40/40 [=====] - 23s 567ms/step - loss: 1.8612 -
OUTPUT_loss: 0.6009 - Softmax_1_loss: 0.6404 - Softmax_2_loss: 0.6199 -
OUTPUT_accuracy: 0.7732 - Softmax_1_accuracy: 0.7625 - Softmax_2_accuracy:
0.7673 - val_loss: 2.5650 - val_OUTPUT_loss: 0.8445 - val_Softmax_1_loss: 0.8687
- val_Softmax_2_loss: 0.8518 - val_OUTPUT_accuracy: 0.7030 -
val_Softmax_1_accuracy: 0.6865 - val_Softmax_2_accuracy: 0.6900

Epoch 00045: LearningRateScheduler reducing learning rate to 0.008153726976.

Epoch 45/50

40/40 [=====] - 23s 567ms/step - loss: 2.1314 -
OUTPUT_loss: 0.6967 - Softmax_1_loss: 0.7288 - Softmax_2_loss: 0.7059 -
OUTPUT_accuracy: 0.7433 - Softmax_1_accuracy: 0.7344 - Softmax_2_accuracy:
0.7378 - val_loss: 3.1333 - val_OUTPUT_loss: 1.0662 - val_Softmax_1_loss: 1.0009
- val_Softmax_2_loss: 1.0662 - val_OUTPUT_accuracy: 0.6400 -
val_Softmax_1_accuracy: 0.6555 - val_Softmax_2_accuracy: 0.6425

Epoch 00046: LearningRateScheduler reducing learning rate to 0.008153726976.

Epoch 46/50

40/40 [=====] - 23s 567ms/step - loss: 2.0166 -
OUTPUT_loss: 0.6545 - Softmax_1_loss: 0.6932 - Softmax_2_loss: 0.6689 -
OUTPUT_accuracy: 0.7584 - Softmax_1_accuracy: 0.7482 - Softmax_2_accuracy:
0.7562 - val_loss: 2.8666 - val_OUTPUT_loss: 0.9801 - val_Softmax_1_loss: 0.9215
- val_Softmax_2_loss: 0.9650 - val_OUTPUT_accuracy: 0.6630 -
val_Softmax_1_accuracy: 0.6695 - val_Softmax_2_accuracy: 0.6640

Epoch 00047: LearningRateScheduler reducing learning rate to 0.008153726976.

Epoch 47/50

40/40 [=====] - 23s 570ms/step - loss: 1.7691 -
OUTPUT_loss: 0.5698 - Softmax_1_loss: 0.6162 - Softmax_2_loss: 0.5832 -
OUTPUT_accuracy: 0.7842 - Softmax_1_accuracy: 0.7698 - Softmax_2_accuracy:
0.7829 - val_loss: 2.6169 - val_OUTPUT_loss: 0.8805 - val_Softmax_1_loss: 0.8646
- val_Softmax_2_loss: 0.8717 - val_OUTPUT_accuracy: 0.7000 -
val_Softmax_1_accuracy: 0.6895 - val_Softmax_2_accuracy: 0.6980

Epoch 00048: LearningRateScheduler reducing learning rate to
0.007827577896959998.

Epoch 48/50

40/40 [=====] - 23s 568ms/step - loss: 1.7475 -
OUTPUT_loss: 0.5602 - Softmax_1_loss: 0.6086 - Softmax_2_loss: 0.5787 -
OUTPUT_accuracy: 0.7921 - Softmax_1_accuracy: 0.7776 - Softmax_2_accuracy:
0.7872 - val_loss: 2.8409 - val_OUTPUT_loss: 0.9324 - val_Softmax_1_loss: 0.9569
- val_Softmax_2_loss: 0.9517 - val_OUTPUT_accuracy: 0.6890 -
val_Softmax_1_accuracy: 0.6690 - val_Softmax_2_accuracy: 0.6820

```
Epoch 00049: LearningRateScheduler reducing learning rate to
0.007827577896959998.
Epoch 49/50
40/40 [=====] - 23s 567ms/step - loss: 1.9356 -
OUTPUT_loss: 0.6241 - Softmax_1_loss: 0.6667 - Softmax_2_loss: 0.6447 -
OUTPUT_accuracy: 0.7663 - Softmax_1_accuracy: 0.7556 - Softmax_2_accuracy:
0.7609 - val_loss: 3.0808 - val_OUTPUT_loss: 1.0472 - val_Softmax_1_loss: 0.9758
- val_Softmax_2_loss: 1.0578 - val_OUTPUT_accuracy: 0.6675 -
val_Softmax_1_accuracy: 0.6690 - val_Softmax_2_accuracy: 0.6595
```

```
Epoch 00050: LearningRateScheduler reducing learning rate to
0.007827577896959998.
Epoch 50/50
40/40 [=====] - 23s 567ms/step - loss: 1.8910 -
OUTPUT_loss: 0.6181 - Softmax_1_loss: 0.6414 - Softmax_2_loss: 0.6315 -
OUTPUT_accuracy: 0.7703 - Softmax_1_accuracy: 0.7609 - Softmax_2_accuracy:
0.7657 - val_loss: 2.6695 - val_OUTPUT_loss: 0.8794 - val_Softmax_1_loss: 0.8799
- val_Softmax_2_loss: 0.9102 - val_OUTPUT_accuracy: 0.6925 -
val_Softmax_1_accuracy: 0.6975 - val_Softmax_2_accuracy: 0.6885
```

```
[17]: model.save('googlenet_n128_colour.h5')
```

```
[20]: from keras.models import load_model
model=load_model('googlenet_n128_colour.h5')
```

```
[24]: score = model.evaluate(test_images,test_labels,verbose=0)
print('Test Loss - ',score[3])
print('Test Accuracy - ',score[5])
```

```
Test Loss - 0.9102452397346497
Test Accuracy - 0.6974999904632568
```

0.11 ## Residual Module

- Residual Model solves the problem of vanishing gradient. The idea is to add a shortcut layer for the flow of information from one layer to another which improves the performance of the model.
- Let $H(x)$ be the underlying mapping to be fit by a deep network, where x is the input to the first layer and the output will be $F(x)$. The residual function is $H(x)-x$ and after approximation we achieve $F(x)+x$. So we add the input x to the output function before feeding it to the activation layer.
- Ref: <https://machinelearningmastery.com/how-to-implement-major-architecture-innovations-f>

```
[11]: #Residual Block :- 2 Conv layer, 64 : 3x3 filters, Relu activation function
def residual_module(layer_in, n_filters):
    merge_input = layer_in
```



```

    # check if the number of filters needs to be increase, assumes channels
    ↪ last format
    if layer_in.shape[-1] != n_filters:
        merge_input = Conv2D(n_filters, (1,1), padding='same',
    ↪ activation='relu', kernel_initializer='he_normal')(layer_in)
        # conv1
        conv1 = Conv2D(n_filters, (3,3), padding='same', activation='relu',
    ↪ kernel_initializer='he_normal')(layer_in)
        # conv2
        conv2 = Conv2D(n_filters, (3,3), padding='same', activation='linear',
    ↪ kernel_initializer='he_normal')(conv1)

    # add filters, assumes filters/channels last
    layer_out = add([conv2, merge_input])
    # activation function
    layer_out = Activation('relu')(layer_out)
    return layer_out

```

0.12 ResNet using residual module

```

[12]: # define model input
visible = Input(shape=(128,128,3))

layer = Conv2D(32, (3,3), padding='same', activation='relu',
    ↪ kernel_initializer='he_normal')(visible) #reduced the filtersize from 32 to 8

#Stacks of 4 residual model
layer = residual_module(visible, 64) #First stack
layer = residual_module(layer, 128) #Second stack
layer = residual_module(layer, 256) #Third stack
layer = residual_module(layer, 512) #Fourth stack

layer = AveragePooling2D((5,5),strides=5)(layer) #Average Pooling layer

layer = Flatten()(layer)

layer = Dense(200, activation='relu')(layer) #FC layer with 200
layer = Dense(len(classes), activation='softmax', name='OUTPUT')(layer) #FC
    ↪ layer with 8 classes

# create model
model = Model(visible,layer)

# summarize model
model.summary()

```

```
# plot model architecture
plot_model(model, show_shapes=True, to_file='residual_module.png')
```

Model: "functional_1"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 128, 128, 3) 0		
conv2d_2 (Conv2D)	(None, 128, 128, 64) 1792		input_1[0][0]
conv2d_3 (Conv2D)	(None, 128, 128, 64) 36928		conv2d_2[0][0]
conv2d_1 (Conv2D)	(None, 128, 128, 64) 256		input_1[0][0]
add (Add)	(None, 128, 128, 64) 0		conv2d_3[0][0] conv2d_1[0][0]
activation (Activation)	(None, 128, 128, 64) 0		add[0][0]
conv2d_5 (Conv2D) activation[0][0]	(None, 128, 128, 128) 73856		
conv2d_6 (Conv2D)	(None, 128, 128, 128) 147584		conv2d_5[0][0]
conv2d_4 (Conv2D) activation[0][0]	(None, 128, 128, 128) 8320		
add_1 (Add)	(None, 128, 128, 128) 0		conv2d_6[0][0] conv2d_4[0][0]
activation_1 (Activation)	(None, 128, 128, 128) 0		add_1[0][0]
conv2d_8 (Conv2D)	(None, 128, 128, 256) 295168		

```

activation_1[0][0]
-----
conv2d_9 (Conv2D)          (None, 128, 128, 256 590080      conv2d_8[0][0]
-----
conv2d_7 (Conv2D)          (None, 128, 128, 256 33024
activation_1[0][0]
-----
add_2 (Add)                (None, 128, 128, 256 0          conv2d_9[0][0]
                                conv2d_7[0][0]
-----
activation_2 (Activation)   (None, 128, 128, 256 0          add_2[0][0]
-----
conv2d_11 (Conv2D)         (None, 128, 128, 512 1180160
activation_2[0][0]
-----
conv2d_12 (Conv2D)         (None, 128, 128, 512 2359808      conv2d_11[0][0]
-----
conv2d_10 (Conv2D)         (None, 128, 128, 512 131584
activation_2[0][0]
-----
add_3 (Add)                (None, 128, 128, 512 0          conv2d_12[0][0]
                                conv2d_10[0][0]
-----
activation_3 (Activation)   (None, 128, 128, 512 0          add_3[0][0]
-----
average_pooling2d (AveragePooli (None, 25, 25, 512) 0
activation_3[0][0]
-----
flatten (Flatten)         (None, 320000)          0
average_pooling2d[0][0]
-----
dense (Dense)              (None, 200)              64000200      flatten[0][0]
-----
OUTPUT (Dense)             (None, 8)                1608          dense[0][0]
=====

```

```
=====
```

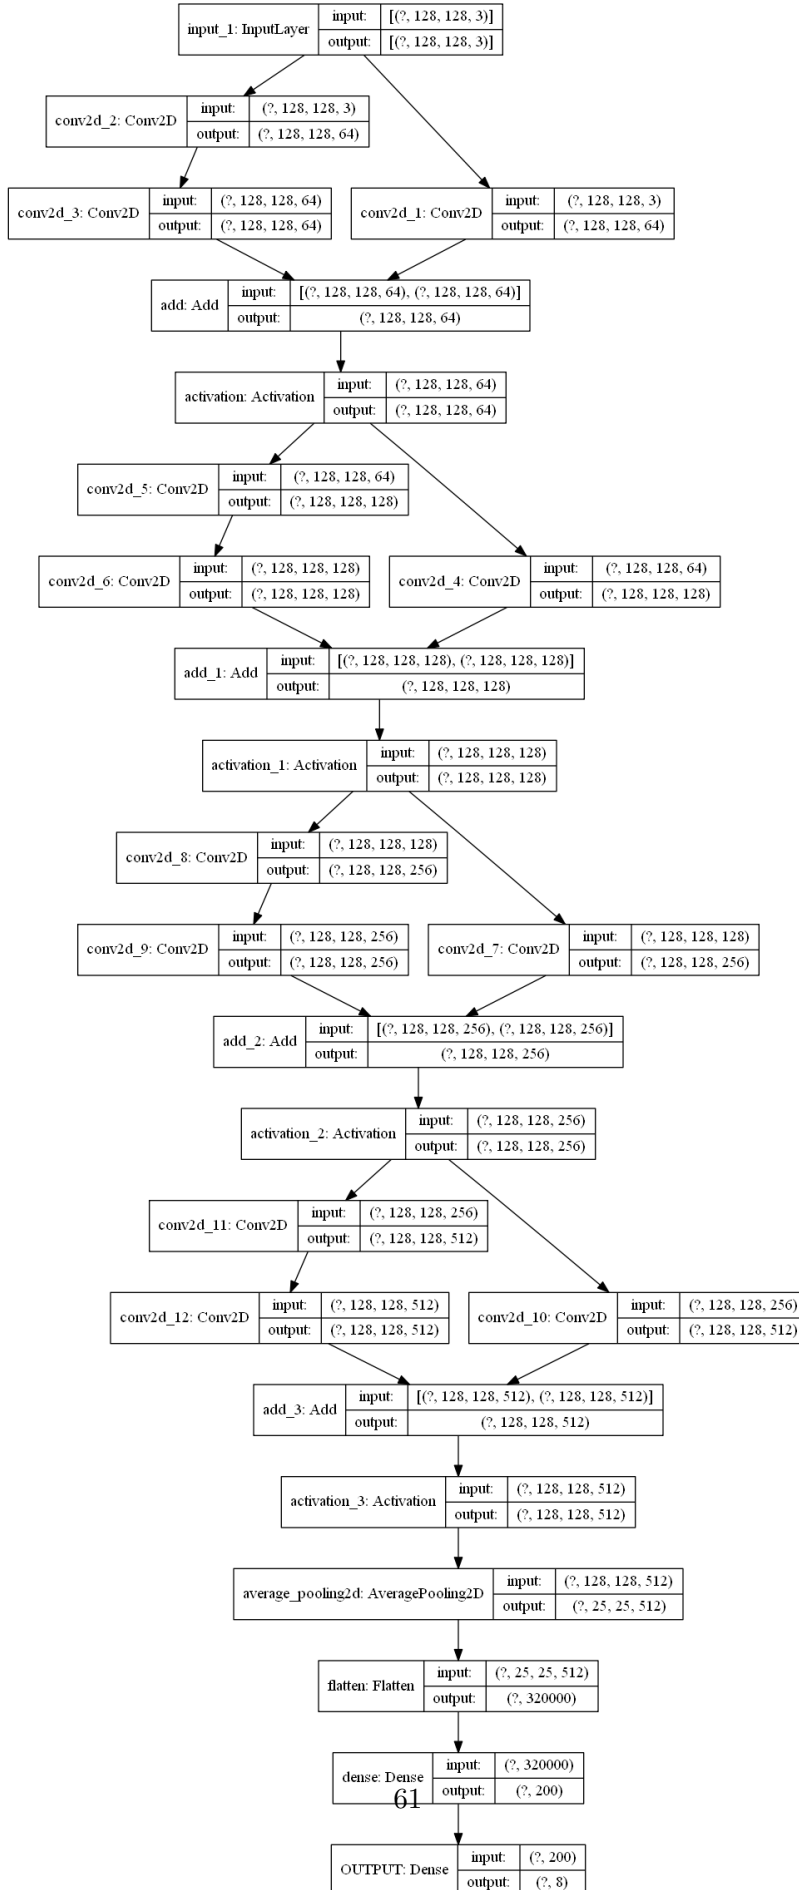
```
Total params: 68,860,368
```

```
Trainable params: 68,860,368
```

```
Non-trainable params: 0
```

```
-----  
-----
```

```
[12]:
```



0.12.1 Model Summary

- We see that we have 68,860,368 trainable parameters. The number of parameters depend on the size of the filter and also the total number of filters used. As this uses a deep dense layer we see such a huge number of parameters.

```
[ ]: #Stochastic Gradient Descent optimizer
opt = SGD(lr=0.01,momentum=0.9,decay=0.01)
from keras.callbacks import EarlyStopping

#Early stopping with a patience of 4 on validation loss
early = EarlyStopping(monitor='val_loss', min_delta=0, patience=4,verbose=0,
    ↪mode='auto')

model.
    ↪compile(optimizer=opt,loss='categorical_crossentropy',metrics=['accuracy'])
```

```
[14]: #Training our 8 class 1500 images per class training dataset on ResNet
start_time=t.time()
history = model.fit(train_images, train_labels,
                    validation_data=(test_images, test_labels),
                    epochs=30, batch_size=4, callbacks=[early])
print("\n Training Time: {} seconds".format(t.time()-start_time))
```

Epoch 1/30

2/2500 [...] - ETA: 5:35 - loss: 14.6347 - accuracy: 0.1250WARNING:tensorflow:Callbacks method `on_train_batch_end` is slow compared to the batch time (batch time: 0.0030s vs `on_train_batch_end` time: 0.2643s). Check your callbacks.

2500/2500 [=====] - ETA: 0s - loss: 2.3877 - accuracy: 0.1193WARNING:tensorflow:Callbacks method `on_test_batch_end` is slow compared to the batch time (batch time: 0.0020s vs `on_test_batch_end` time: 0.1023s). Check your callbacks.

2500/2500 [=====] - 717s 287ms/step - loss: 2.3877 - accuracy: 0.1193 - val_loss: 2.0801 - val_accuracy: 0.1195

Epoch 2/30

2500/2500 [=====] - 717s 287ms/step - loss: 2.0796 - accuracy: 0.1219 - val_loss: 2.0801 - val_accuracy: 0.1150

Epoch 3/30

2500/2500 [=====] - 717s 287ms/step - loss: 2.0794 - accuracy: 0.1278 - val_loss: 2.0797 - val_accuracy: 0.1155

Epoch 4/30

2500/2500 [=====] - 717s 287ms/step - loss: 2.0789 - accuracy: 0.1293 - val_loss: 2.0787 - val_accuracy: 0.1245

```

Epoch 5/30
2500/2500 [=====] - 717s 287ms/step - loss: 2.0752 -
accuracy: 0.1377 - val_loss: 2.0536 - val_accuracy: 0.1830
Epoch 6/30
2500/2500 [=====] - 718s 287ms/step - loss: 2.0437 -
accuracy: 0.1833 - val_loss: 1.9830 - val_accuracy: 0.2025
Epoch 7/30
2500/2500 [=====] - 718s 287ms/step - loss: 1.9782 -
accuracy: 0.2276 - val_loss: 1.9755 - val_accuracy: 0.2270
Epoch 8/30
2500/2500 [=====] - 718s 287ms/step - loss: 1.9353 -
accuracy: 0.2563 - val_loss: 1.9223 - val_accuracy: 0.2710
Epoch 9/30
2500/2500 [=====] - 718s 287ms/step - loss: 1.9047 -
accuracy: 0.2667 - val_loss: 1.9208 - val_accuracy: 0.2605
Epoch 10/30
2500/2500 [=====] - 718s 287ms/step - loss: 1.8714 -
accuracy: 0.2975 - val_loss: 1.9086 - val_accuracy: 0.2480
Epoch 11/30
2500/2500 [=====] - 717s 287ms/step - loss: 1.8472 -
accuracy: 0.3047 - val_loss: 1.8980 - val_accuracy: 0.2680
Epoch 12/30
2500/2500 [=====] - 717s 287ms/step - loss: 1.8156 -
accuracy: 0.3185 - val_loss: 1.8935 - val_accuracy: 0.2610
Epoch 13/30
2500/2500 [=====] - 717s 287ms/step - loss: 1.7931 -
accuracy: 0.3342 - val_loss: 1.8693 - val_accuracy: 0.2950
Epoch 14/30
2500/2500 [=====] - 717s 287ms/step - loss: 1.7706 -
accuracy: 0.3450 - val_loss: 1.8786 - val_accuracy: 0.2915
Epoch 15/30
2500/2500 [=====] - 717s 287ms/step - loss: 1.7488 -
accuracy: 0.3565 - val_loss: 1.8801 - val_accuracy: 0.2750
Epoch 16/30
2500/2500 [=====] - 717s 287ms/step - loss: 1.7208 -
accuracy: 0.3655 - val_loss: 2.0175 - val_accuracy: 0.2675
Epoch 17/30
2500/2500 [=====] - 717s 287ms/step - loss: 1.7037 -
accuracy: 0.3749 - val_loss: 1.8811 - val_accuracy: 0.2950

```

Training Time: 12207.317234992981 seconds

```

[17]: #Saving the model
model.save('resnet_n128_colour.h5')

```

0.13 ## Combined Model 1

- Inspired from AlexNet and Inception module

```
[36]: input_layer = Input(shape=(n, n, 3))
m = Conv2D(96, kernel_size=(11,11), strides=(4,4), padding="valid",
↪activation="relu")(input_layer)
m = MaxPooling2D(pool_size=(3,3), strides=(2,2))(m)
m = Conv2D(256, kernel_size=(5,5), padding="same", strides=(1,1),
↪activation="relu")(m)
m = MaxPooling2D(pool_size=(3,3), strides=(2,2))(m)
m = inception_module(m, 160, 112, 224, 24, 64, 64)
m = inception_module(m, 128, 128, 256, 24, 64, 64)
m = Conv2D(384, kernel_size=(3,3), padding="same", strides=(1,1),
↪activation="relu")(m)
m = Conv2D(256, kernel_size=(3,3), padding="same", strides=(1,1),
↪activation="relu")(m)
m = MaxPooling2D(pool_size=(3,3), strides=(2,2))(m)
m = Flatten()(m)
m = Dense(4096, activation="relu")(m)
m = Dense(units=len(classes), activation="softmax", name = "OUTPUT")(m)

model4 = Model(input_layer, m, name='Combination_1')

model4.summary()
plot_model(model4)
```

Model: "Combination_1"

Layer (type)	Output Shape	Param #	Connected to
input_9 (InputLayer)	[(None, 128, 128, 3)]	0	
conv2d_88 (Conv2D)	(None, 30, 30, 96)	34944	input_9[0][0]
max_pooling2d_21 (MaxPooling2D)	(None, 14, 14, 96)	0	conv2d_88[0][0]
conv2d_89 (Conv2D)	(None, 14, 14, 256)	614656	max_pooling2d_21[0][0]
max_pooling2d_22 (MaxPooling2D)	(None, 6, 6, 256)	0	conv2d_89[0][0]
conv2d_91 (Conv2D)	(None, 6, 6, 112)	28784	max_pooling2d_22[0][0]


```

-----
conv2d_93 (Conv2D) (None, 6, 6, 24) 6168
max_pooling2d_22[0][0]
-----

max_pooling2d_23 (MaxPooling2D) (None, 6, 6, 256) 0
max_pooling2d_22[0][0]
-----

conv2d_90 (Conv2D) (None, 6, 6, 160) 41120
max_pooling2d_22[0][0]
-----

conv2d_92 (Conv2D) (None, 6, 6, 224) 226016 conv2d_91[0][0]
-----

conv2d_94 (Conv2D) (None, 6, 6, 64) 38464 conv2d_93[0][0]
-----

conv2d_95 (Conv2D) (None, 6, 6, 64) 16448
max_pooling2d_23[0][0]
-----

concatenate_12 (Concatenate) (None, 6, 6, 512) 0 conv2d_90[0][0]
conv2d_92[0][0]
conv2d_94[0][0]
conv2d_95[0][0]
-----

conv2d_97 (Conv2D) (None, 6, 6, 128) 65664
concatenate_12[0][0]
-----

conv2d_99 (Conv2D) (None, 6, 6, 24) 12312
concatenate_12[0][0]
-----

max_pooling2d_24 (MaxPooling2D) (None, 6, 6, 512) 0
concatenate_12[0][0]
-----

conv2d_96 (Conv2D) (None, 6, 6, 128) 65664
concatenate_12[0][0]
-----

conv2d_98 (Conv2D) (None, 6, 6, 256) 295168 conv2d_97[0][0]
-----

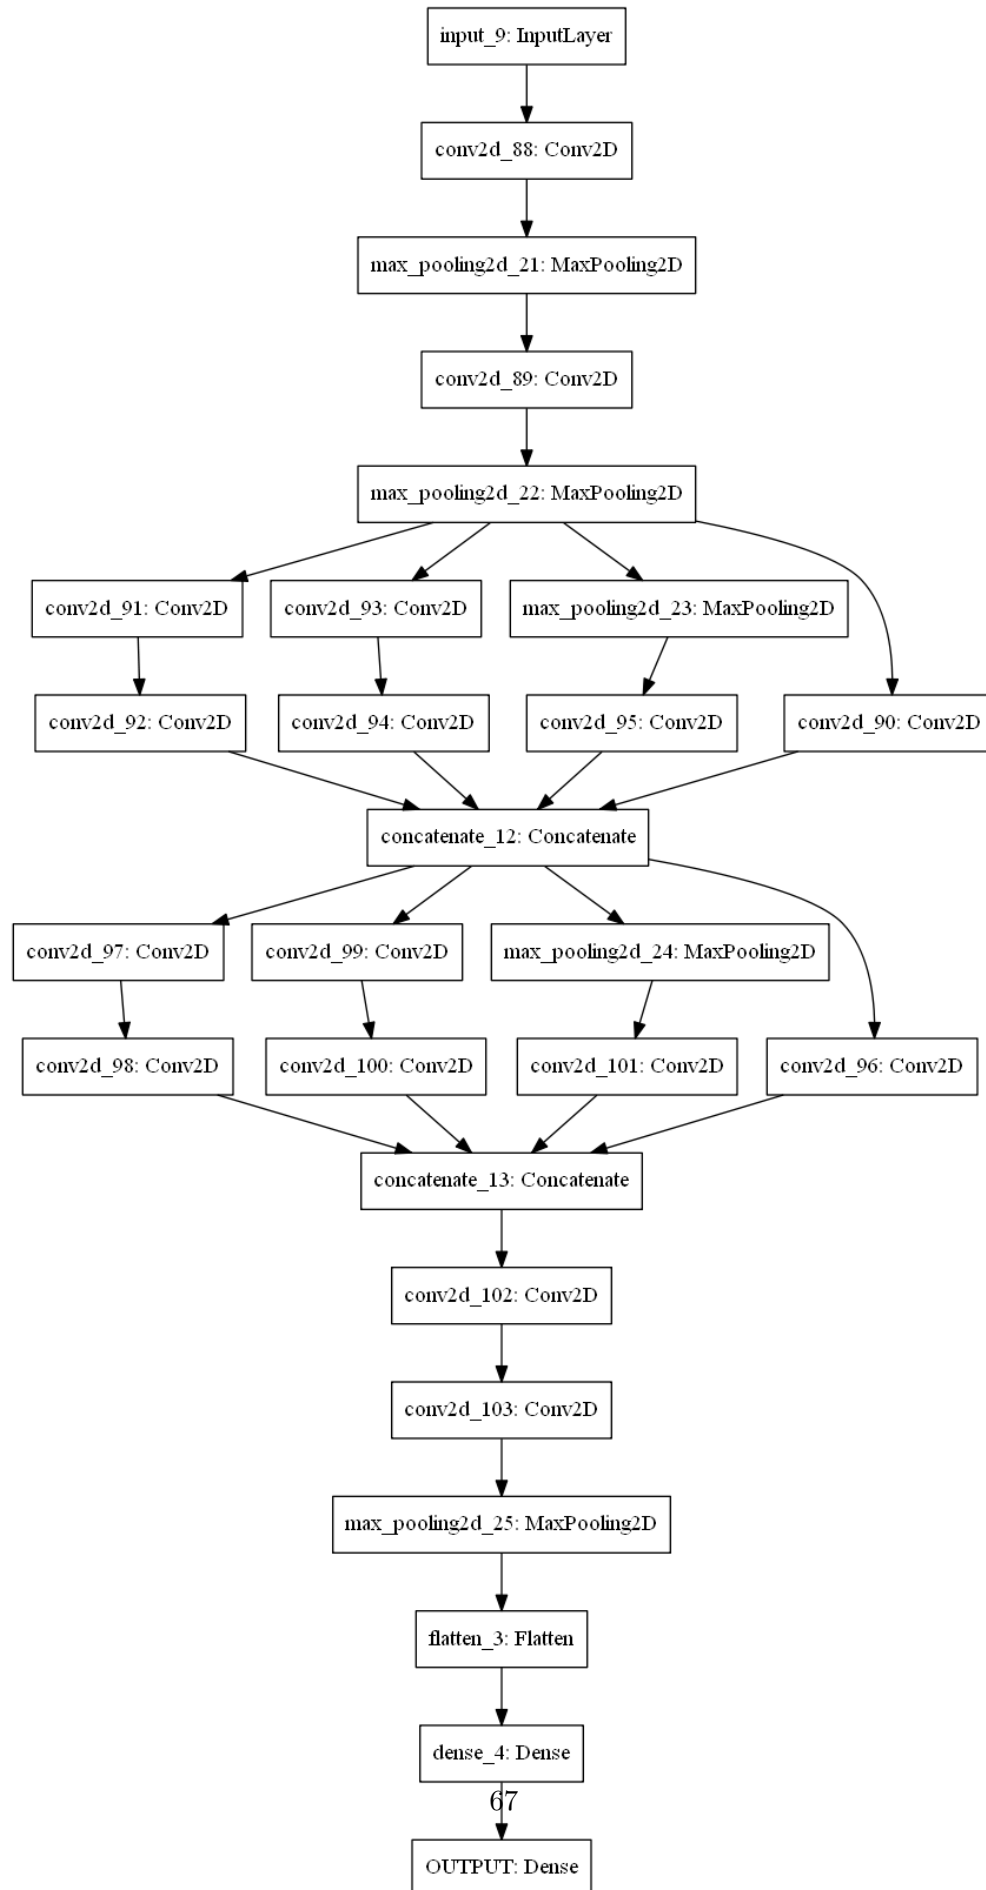
```

```

-----
conv2d_100 (Conv2D)          (None, 6, 6, 64)      38464      conv2d_99[0][0]
-----
conv2d_101 (Conv2D)          (None, 6, 6, 64)      32832
max_pooling2d_24[0][0]
-----
concatenate_13 (Concatenate) (None, 6, 6, 512)      0          conv2d_96[0][0]
                                         conv2d_98[0][0]
conv2d_100[0][0]
conv2d_101[0][0]
-----
conv2d_102 (Conv2D)          (None, 6, 6, 384)     1769856
concatenate_13[0][0]
-----
conv2d_103 (Conv2D)          (None, 6, 6, 256)     884992
conv2d_102[0][0]
-----
max_pooling2d_25 (MaxPooling2D) (None, 2, 2, 256)      0
conv2d_103[0][0]
-----
flatten_3 (Flatten)          (None, 1024)          0
max_pooling2d_25[0][0]
-----
dense_4 (Dense)              (None, 4096)          4198400     flatten_3[0][0]
-----
OUTPUT (Dense)              (None, 8)             32776       dense_4[0][0]
=====
Total params: 8,402,728
Trainable params: 8,402,728
Non-trainable params: 0
-----

```

[36]:



```
[11]: opt = SGD(lr=0.01,momentum=0.9,decay=0.01)

[14]: from keras.callbacks import ReduceLROnPlateau, EarlyStopping
early = EarlyStopping(monitor='val_loss', min_delta=0, patience=10,verbose=0,
↳mode='auto')

[40]: model4.
↳compile(optimizer=opt,loss='categorical_crossentropy',metrics=['accuracy'])

[43]: start_time=t.time()
history4 = model4.fit(train_images, train_labels,
validation_data=(test_images, test_labels),
epochs=100, batch_size=64,callbacks=[early])
print("\n Training Time: {} seconds".format(t.time()-start_time))
```

```
Epoch 1/100
157/157 [=====] - 14s 91ms/step - loss: 2.0319 -
accuracy: 0.1849 - val_loss: 1.9150 - val_accuracy: 0.2515
Epoch 2/100
157/157 [=====] - 10s 65ms/step - loss: 1.8166 -
accuracy: 0.2834 - val_loss: 2.1114 - val_accuracy: 0.2200
Epoch 3/100
157/157 [=====] - 11s 69ms/step - loss: 1.6566 -
accuracy: 0.3701 - val_loss: 1.4968 - val_accuracy: 0.4470
Epoch 4/100
157/157 [=====] - 10s 65ms/step - loss: 1.4466 -
accuracy: 0.4558 - val_loss: 1.3831 - val_accuracy: 0.4925
Epoch 5/100
157/157 [=====] - 11s 68ms/step - loss: 1.3376 -
accuracy: 0.4998 - val_loss: 1.2881 - val_accuracy: 0.5420
Epoch 6/100
157/157 [=====] - 11s 68ms/step - loss: 1.2407 -
accuracy: 0.5417 - val_loss: 1.2111 - val_accuracy: 0.5590
Epoch 7/100
157/157 [=====] - 11s 68ms/step - loss: 1.1865 -
accuracy: 0.5605 - val_loss: 1.1884 - val_accuracy: 0.5660
Epoch 8/100
157/157 [=====] - 11s 68ms/step - loss: 1.1564 -
accuracy: 0.5701 - val_loss: 1.1737 - val_accuracy: 0.5690
Epoch 9/100
157/157 [=====] - 11s 71ms/step - loss: 1.1136 -
accuracy: 0.5865 - val_loss: 1.2325 - val_accuracy: 0.5505
Epoch 10/100
157/157 [=====] - 11s 69ms/step - loss: 1.1058 -
accuracy: 0.5896 - val_loss: 1.1485 - val_accuracy: 0.5940
```

Epoch 11/100
157/157 [=====] - 10s 66ms/step - loss: 1.0583 - accuracy: 0.6114 - val_loss: 1.0973 - val_accuracy: 0.6095

Epoch 12/100
157/157 [=====] - 10s 67ms/step - loss: 1.0229 - accuracy: 0.6182 - val_loss: 1.1073 - val_accuracy: 0.5930

Epoch 13/100
157/157 [=====] - 11s 71ms/step - loss: 1.0037 - accuracy: 0.6275 - val_loss: 1.0658 - val_accuracy: 0.6155

Epoch 14/100
157/157 [=====] - 11s 70ms/step - loss: 0.9649 - accuracy: 0.6417 - val_loss: 1.1621 - val_accuracy: 0.5835

Epoch 15/100
157/157 [=====] - 10s 66ms/step - loss: 0.9540 - accuracy: 0.6463 - val_loss: 1.0625 - val_accuracy: 0.6260

Epoch 16/100
157/157 [=====] - 11s 69ms/step - loss: 0.9422 - accuracy: 0.6508 - val_loss: 1.0335 - val_accuracy: 0.6260

Epoch 17/100
157/157 [=====] - 11s 69ms/step - loss: 0.9237 - accuracy: 0.6592 - val_loss: 1.0348 - val_accuracy: 0.6170

Epoch 18/100
157/157 [=====] - 11s 71ms/step - loss: 0.8915 - accuracy: 0.6688 - val_loss: 1.0017 - val_accuracy: 0.6400

Epoch 19/100
157/157 [=====] - 10s 66ms/step - loss: 0.8826 - accuracy: 0.6713 - val_loss: 1.0509 - val_accuracy: 0.6180

Epoch 20/100
157/157 [=====] - 12s 74ms/step - loss: 0.8593 - accuracy: 0.6843 - val_loss: 0.9898 - val_accuracy: 0.6365

Epoch 21/100
157/157 [=====] - 12s 74ms/step - loss: 0.8292 - accuracy: 0.6939 - val_loss: 1.0013 - val_accuracy: 0.6325

Epoch 22/100
157/157 [=====] - 11s 70ms/step - loss: 0.8298 - accuracy: 0.6940 - val_loss: 1.0019 - val_accuracy: 0.6345

Epoch 23/100
157/157 [=====] - 11s 70ms/step - loss: 0.8115 - accuracy: 0.6973 - val_loss: 0.9976 - val_accuracy: 0.6375

Epoch 24/100
157/157 [=====] - 11s 68ms/step - loss: 0.8031 - accuracy: 0.7030 - val_loss: 0.9980 - val_accuracy: 0.6400

Epoch 25/100
157/157 [=====] - 11s 72ms/step - loss: 0.7813 - accuracy: 0.7111 - val_loss: 0.9746 - val_accuracy: 0.6450

Epoch 26/100
157/157 [=====] - 11s 68ms/step - loss: 0.7662 - accuracy: 0.7116 - val_loss: 0.9414 - val_accuracy: 0.6570

Epoch 27/100
157/157 [=====] - 11s 70ms/step - loss: 0.7462 -
accuracy: 0.7259 - val_loss: 0.9242 - val_accuracy: 0.6675
Epoch 28/100
157/157 [=====] - 11s 72ms/step - loss: 0.7371 -
accuracy: 0.7297 - val_loss: 0.9585 - val_accuracy: 0.6585
Epoch 29/100
157/157 [=====] - 12s 75ms/step - loss: 0.7396 -
accuracy: 0.7266 - val_loss: 0.9479 - val_accuracy: 0.6555
Epoch 30/100
157/157 [=====] - 11s 69ms/step - loss: 0.7203 -
accuracy: 0.7340 - val_loss: 0.9345 - val_accuracy: 0.6595
Epoch 31/100
157/157 [=====] - 11s 72ms/step - loss: 0.7080 -
accuracy: 0.7402 - val_loss: 0.9526 - val_accuracy: 0.6570
Epoch 32/100
157/157 [=====] - 12s 74ms/step - loss: 0.7067 -
accuracy: 0.7376 - val_loss: 0.9042 - val_accuracy: 0.6685
Epoch 33/100
157/157 [=====] - 11s 71ms/step - loss: 0.6818 -
accuracy: 0.7512 - val_loss: 0.9598 - val_accuracy: 0.6550
Epoch 34/100
157/157 [=====] - 11s 69ms/step - loss: 0.6743 -
accuracy: 0.7502 - val_loss: 0.9002 - val_accuracy: 0.6825
Epoch 35/100
157/157 [=====] - 11s 71ms/step - loss: 0.6635 -
accuracy: 0.7558 - val_loss: 0.9719 - val_accuracy: 0.6625
Epoch 36/100
157/157 [=====] - 11s 72ms/step - loss: 0.6629 -
accuracy: 0.7511 - val_loss: 0.8976 - val_accuracy: 0.6775
Epoch 37/100
157/157 [=====] - 12s 75ms/step - loss: 0.6455 -
accuracy: 0.7631 - val_loss: 0.9525 - val_accuracy: 0.6620
Epoch 38/100
157/157 [=====] - 11s 73ms/step - loss: 0.6393 -
accuracy: 0.7625 - val_loss: 0.9135 - val_accuracy: 0.6855
Epoch 39/100
157/157 [=====] - 11s 68ms/step - loss: 0.6202 -
accuracy: 0.7720 - val_loss: 0.9185 - val_accuracy: 0.6775
Epoch 40/100
157/157 [=====] - 11s 72ms/step - loss: 0.6165 -
accuracy: 0.7722 - val_loss: 0.9052 - val_accuracy: 0.6740
Epoch 41/100
157/157 [=====] - 12s 74ms/step - loss: 0.5972 -
accuracy: 0.7806 - val_loss: 0.9384 - val_accuracy: 0.6620
Epoch 42/100
157/157 [=====] - 10s 67ms/step - loss: 0.5832 -
accuracy: 0.7865 - val_loss: 0.9103 - val_accuracy: 0.6785



Epoch 43/100
157/157 [=====] - 11s 73ms/step - loss: 0.5849 -
accuracy: 0.7851 - val_loss: 0.9672 - val_accuracy: 0.6615
Epoch 44/100
157/157 [=====] - 12s 74ms/step - loss: 0.5596 -
accuracy: 0.7908 - val_loss: 0.8924 - val_accuracy: 0.6800
Epoch 45/100
157/157 [=====] - 11s 71ms/step - loss: 0.5594 -
accuracy: 0.7921 - val_loss: 0.8860 - val_accuracy: 0.6860
Epoch 46/100
157/157 [=====] - 11s 71ms/step - loss: 0.5502 -
accuracy: 0.7936 - val_loss: 0.9758 - val_accuracy: 0.6650
Epoch 47/100
157/157 [=====] - 11s 68ms/step - loss: 0.5448 -
accuracy: 0.7967 - val_loss: 0.8759 - val_accuracy: 0.6925
Epoch 48/100
157/157 [=====] - 11s 68ms/step - loss: 0.5396 -
accuracy: 0.7995 - val_loss: 0.9174 - val_accuracy: 0.6705
Epoch 49/100
157/157 [=====] - 11s 68ms/step - loss: 0.5306 -
accuracy: 0.8056 - val_loss: 0.9079 - val_accuracy: 0.6850
Epoch 50/100
157/157 [=====] - 12s 75ms/step - loss: 0.5061 -
accuracy: 0.8126 - val_loss: 0.9557 - val_accuracy: 0.6785
Epoch 51/100
157/157 [=====] - 12s 74ms/step - loss: 0.5110 -
accuracy: 0.8112 - val_loss: 0.9754 - val_accuracy: 0.6730
Epoch 52/100
157/157 [=====] - 11s 69ms/step - loss: 0.4931 -
accuracy: 0.8166 - val_loss: 0.9103 - val_accuracy: 0.6830
Epoch 53/100
157/157 [=====] - 11s 71ms/step - loss: 0.4922 -
accuracy: 0.8160 - val_loss: 0.9778 - val_accuracy: 0.6665
Epoch 54/100
157/157 [=====] - 11s 71ms/step - loss: 0.4812 -
accuracy: 0.8233 - val_loss: 0.8953 - val_accuracy: 0.6875
Epoch 55/100
157/157 [=====] - 11s 69ms/step - loss: 0.4709 -
accuracy: 0.8271 - val_loss: 0.9242 - val_accuracy: 0.6810
Epoch 56/100
157/157 [=====] - 11s 70ms/step - loss: 0.4488 -
accuracy: 0.8367 - val_loss: 0.9007 - val_accuracy: 0.6915
Epoch 57/100
157/157 [=====] - 11s 69ms/step - loss: 0.4659 -
accuracy: 0.8275 - val_loss: 1.0141 - val_accuracy: 0.6745

Training Time: 651.438455581665 seconds

```
[45]: score4 = model4.evaluate(test_images, test_labels, verbose=0)
print('Test Loss - ', score4[0])
print('Test Accuracy - ', score4[1])
graphs(history4, score4)
```

Test Loss - 1.0140718221664429

Test Accuracy - 0.6744999885559082



Accuracy of the model: 67.45%

0.14 ## Combination model 2

- Reduced the overfitting in model 1
- Added another inception module

```
[15]: #combination 2
input_layer = Input(shape=(n, n, 3))
m = Conv2D(96,kernel_size=(11,11),strides=(4,4),kernel_initializer=kernel_init,
↪bias_initializer=bias_init,padding="valid", activation="relu")(input_layer)
m = MaxPooling2D(pool_size=(3,3),strides=(2,2))(m)
m = Conv2D(256, kernel_size=(5,5), padding="same",strides=(1,1),
↪activation="relu")(m)
m = MaxPooling2D(pool_size=(3,3),strides=(2,2))(m)
m = inception_module(m, 160, 112, 224, 24, 64, 64)
m = Dropout(0.3)(m)
m = Conv2D(384, kernel_size=(3,3), padding="same", strides=(1,1),
↪activation="relu")(m)
m = BatchNormalization()(m)
m = Conv2D(256, kernel_size=(3,3), padding="same", strides=(1,1),
↪activation="relu")(m)
m = inception_module(m, 128, 128, 256, 24, 64, 64)
m = Dropout(0.5)(m)
m = inception_module(m, 128, 128, 256, 24, 64, 64)
m = Conv2D(256, kernel_size=(3,3), padding="same", strides=(1,1),
↪activation="relu")(m)
m = MaxPooling2D(pool_size=(3,3),strides=(2,2))(m)
m = Flatten()(m)
m = Dense(4096,activation="relu")(m)
m = Dense(4096,activation="relu")(m)
m = Dense(units=len(classes), activation="softmax", name = "OUTPUT")(m)

model5 = Model(input_layer, m, name='Combination_1')

model5.summary()
plot_model(model5)
```

Model: "Combination_1"

Layer (type)	Output Shape	Param #	Connected to
input_2 (InputLayer)	[(None, 128, 128, 3) 0		
conv2d_23 (Conv2D)	(None, 30, 30, 96)	34944	input_2[0][0]

max_pooling2d_6 (MaxPooling2D)	(None, 14, 14, 96)	0	conv2d_23[0][0]

conv2d_24 (Conv2D)	(None, 14, 14, 256)	614656	
max_pooling2d_6[0][0]			

max_pooling2d_7 (MaxPooling2D)	(None, 6, 6, 256)	0	conv2d_24[0][0]

conv2d_26 (Conv2D)	(None, 6, 6, 112)	28784	
max_pooling2d_7[0][0]			

conv2d_28 (Conv2D)	(None, 6, 6, 24)	6168	
max_pooling2d_7[0][0]			

max_pooling2d_8 (MaxPooling2D)	(None, 6, 6, 256)	0	
max_pooling2d_7[0][0]			

conv2d_25 (Conv2D)	(None, 6, 6, 160)	41120	
max_pooling2d_7[0][0]			

conv2d_27 (Conv2D)	(None, 6, 6, 224)	226016	conv2d_26[0][0]

conv2d_29 (Conv2D)	(None, 6, 6, 64)	38464	conv2d_28[0][0]

conv2d_30 (Conv2D)	(None, 6, 6, 64)	16448	
max_pooling2d_8[0][0]			

concatenate_3 (Concatenate)	(None, 6, 6, 512)	0	conv2d_25[0][0] conv2d_27[0][0] conv2d_29[0][0] conv2d_30[0][0]

dropout_2 (Dropout)	(None, 6, 6, 512)	0	
concatenate_3[0][0]			

conv2d_31 (Conv2D)	(None, 6, 6, 384)	1769856	dropout_2[0][0]

batch_normalization_1 (BatchNor	(None, 6, 6, 384)	1536	conv2d_31[0][0]

conv2d_32 (Conv2D)	(None, 6, 6, 256)	884992	
batch_normalization_1[0][0]			

conv2d_34 (Conv2D)	(None, 6, 6, 128)	32896	conv2d_32[0][0]

conv2d_36 (Conv2D)	(None, 6, 6, 24)	6168	conv2d_32[0][0]

max_pooling2d_9 (MaxPooling2D)	(None, 6, 6, 256)	0	conv2d_32[0][0]

conv2d_33 (Conv2D)	(None, 6, 6, 128)	32896	conv2d_32[0][0]

conv2d_35 (Conv2D)	(None, 6, 6, 256)	295168	conv2d_34[0][0]

conv2d_37 (Conv2D)	(None, 6, 6, 64)	38464	conv2d_36[0][0]

conv2d_38 (Conv2D)	(None, 6, 6, 64)	16448	
max_pooling2d_9[0][0]			

concatenate_4 (Concatenate)	(None, 6, 6, 512)	0	conv2d_33[0][0] conv2d_35[0][0] conv2d_37[0][0] conv2d_38[0][0]

dropout_3 (Dropout)	(None, 6, 6, 512)	0	
concatenate_4[0][0]			

conv2d_40 (Conv2D)	(None, 6, 6, 128)	65664	dropout_3[0][0]

conv2d_42 (Conv2D)	(None, 6, 6, 24)	12312	dropout_3[0][0]

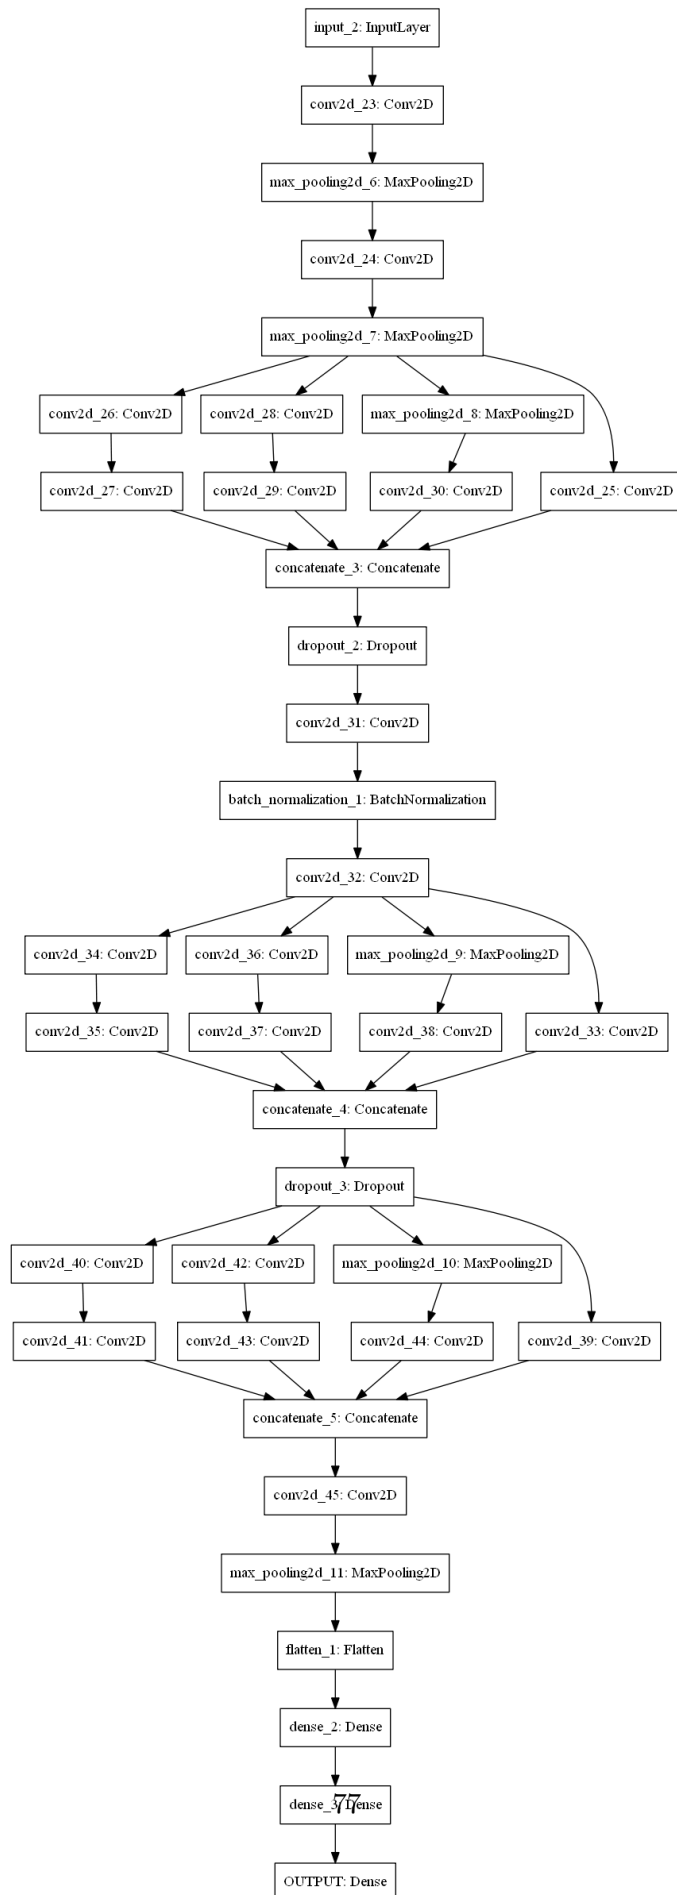
max_pooling2d_10 (MaxPooling2D)	(None, 6, 6, 512)	0	dropout_3[0][0]

```

-----
conv2d_39 (Conv2D)          (None, 6, 6, 128)    65664      dropout_3[0][0]
-----
conv2d_41 (Conv2D)          (None, 6, 6, 256)    295168     conv2d_40[0][0]
-----
conv2d_43 (Conv2D)          (None, 6, 6, 64)     38464      conv2d_42[0][0]
-----
conv2d_44 (Conv2D)          (None, 6, 6, 64)     32832
max_pooling2d_10[0][0]
-----
concatenate_5 (Concatenate) (None, 6, 6, 512)    0           conv2d_39[0][0]
                                           conv2d_41[0][0]
                                           conv2d_43[0][0]
                                           conv2d_44[0][0]
-----
conv2d_45 (Conv2D)          (None, 6, 6, 256)    1179904
concatenate_5[0][0]
-----
max_pooling2d_11 (MaxPooling2D) (None, 2, 2, 256)    0           conv2d_45[0][0]
-----
flatten_1 (Flatten)         (None, 1024)         0
max_pooling2d_11[0][0]
-----
dense_2 (Dense)             (None, 4096)         4198400    flatten_1[0][0]
-----
dense_3 (Dense)             (None, 4096)         16781312   dense_2[0][0]
-----
OUTPUT (Dense)              (None, 8)            32776      dense_3[0][0]
=====
Total params: 26,787,520
Trainable params: 26,786,752
Non-trainable params: 768
-----

```

[15]:



```
[16]: initial_lrate = 0.01
def decay(epoch, steps=100):
    initial_lrate = 0.01
    drop = 0.96
    epochs_drop = 8
    lrate = initial_lrate * math.pow(drop, math.floor((1+epoch)/epochs_drop))
    return lrate

sgd = SGD(lr=initial_lrate, momentum=0.9, nesterov=False)

lr_sc = LearningRateScheduler(decay, verbose=1)

model5.compile(loss='categorical_crossentropy', optimizer=sgd,
               metrics=['accuracy'])

[18]: start_time=t.time()
history5 = model5.fit(train_images, train_labels,
                      validation_data=(test_images, test_labels),
                      epochs=100, batch_size=64, callbacks=[early, lr_sc])
print("\n Training Time: {} seconds".format(t.time()-start_time))
```

Epoch 00001: LearningRateScheduler reducing learning rate to 0.01.

Epoch 1/100

2/157 [...] - ETA: 5s - loss: 2.2022 - accuracy: 0.1562
 WARNING:tensorflow:Callbacks method `on_train_batch_end` is slow compared to the batch time (batch time: 0.0270s vs `on_train_batch_end` time: 0.0409s). Check your callbacks.
 157/157 [=====] - 12s 79ms/step - loss: 2.0067 - accuracy: 0.1898 - val_loss: 2.0169 - val_accuracy: 0.2345

Epoch 00002: LearningRateScheduler reducing learning rate to 0.01.

Epoch 2/100

157/157 [=====] - 11s 70ms/step - loss: 1.7394 - accuracy: 0.3131 - val_loss: 1.7675 - val_accuracy: 0.2940

Epoch 00003: LearningRateScheduler reducing learning rate to 0.01.

Epoch 3/100

157/157 [=====] - 11s 70ms/step - loss: 1.4978 - accuracy: 0.4201 - val_loss: 1.6466 - val_accuracy: 0.3695

Epoch 00004: LearningRateScheduler reducing learning rate to 0.01.

Epoch 4/100

157/157 [=====] - 11s 70ms/step - loss: 1.3442 - accuracy: 0.4932 - val_loss: 1.4556 - val_accuracy: 0.4570

Epoch 00005: LearningRateScheduler reducing learning rate to 0.01.
Epoch 5/100
157/157 [=====] - 11s 71ms/step - loss: 1.2883 -
accuracy: 0.5114 - val_loss: 1.5681 - val_accuracy: 0.4350

Epoch 00006: LearningRateScheduler reducing learning rate to 0.01.
Epoch 6/100
157/157 [=====] - 11s 71ms/step - loss: 1.1929 -
accuracy: 0.5514 - val_loss: 1.3907 - val_accuracy: 0.4630

Epoch 00007: LearningRateScheduler reducing learning rate to 0.01.
Epoch 7/100
157/157 [=====] - 11s 71ms/step - loss: 1.1139 -
accuracy: 0.5863 - val_loss: 2.2350 - val_accuracy: 0.2765

Epoch 00008: LearningRateScheduler reducing learning rate to 0.0096.
Epoch 8/100
157/157 [=====] - 11s 71ms/step - loss: 1.0404 -
accuracy: 0.6119 - val_loss: 1.5471 - val_accuracy: 0.4285

Epoch 00009: LearningRateScheduler reducing learning rate to 0.0096.
Epoch 9/100
157/157 [=====] - 11s 71ms/step - loss: 0.9790 -
accuracy: 0.6355 - val_loss: 1.2087 - val_accuracy: 0.5595

Epoch 00010: LearningRateScheduler reducing learning rate to 0.0096.
Epoch 10/100
157/157 [=====] - 11s 71ms/step - loss: 0.9430 -
accuracy: 0.6522 - val_loss: 1.6276 - val_accuracy: 0.4505

Epoch 00011: LearningRateScheduler reducing learning rate to 0.0096.
Epoch 11/100
157/157 [=====] - 11s 71ms/step - loss: 0.8896 -
accuracy: 0.6676 - val_loss: 1.0411 - val_accuracy: 0.6015

Epoch 00012: LearningRateScheduler reducing learning rate to 0.0096.
Epoch 12/100
157/157 [=====] - 11s 71ms/step - loss: 0.8635 -
accuracy: 0.6810 - val_loss: 1.2039 - val_accuracy: 0.5565

Epoch 00013: LearningRateScheduler reducing learning rate to 0.0096.
Epoch 13/100
157/157 [=====] - 11s 71ms/step - loss: 0.8047 -
accuracy: 0.6974 - val_loss: 1.6535 - val_accuracy: 0.4650

Epoch 00014: LearningRateScheduler reducing learning rate to 0.0096.
Epoch 14/100

157/157 [=====] - 11s 71ms/step - loss: 0.7526 -
accuracy: 0.7118 - val_loss: 1.1366 - val_accuracy: 0.6040

Epoch 00015: LearningRateScheduler reducing learning rate to 0.0096.

Epoch 15/100

157/157 [=====] - 11s 71ms/step - loss: 0.7164 -
accuracy: 0.7280 - val_loss: 1.1746 - val_accuracy: 0.5910

Epoch 00016: LearningRateScheduler reducing learning rate to 0.009216.

Epoch 16/100

157/157 [=====] - 11s 71ms/step - loss: 0.6975 -
accuracy: 0.7419 - val_loss: 1.1037 - val_accuracy: 0.5885

Epoch 00017: LearningRateScheduler reducing learning rate to 0.009216.

Epoch 17/100

157/157 [=====] - 11s 71ms/step - loss: 0.6417 -
accuracy: 0.7574 - val_loss: 1.3874 - val_accuracy: 0.5100

Epoch 00018: LearningRateScheduler reducing learning rate to 0.009216.

Epoch 18/100

157/157 [=====] - 11s 71ms/step - loss: 0.6073 -
accuracy: 0.7733 - val_loss: 1.0211 - val_accuracy: 0.6355

Epoch 00019: LearningRateScheduler reducing learning rate to 0.009216.

Epoch 19/100

157/157 [=====] - 11s 71ms/step - loss: 0.5751 -
accuracy: 0.7810 - val_loss: 1.4679 - val_accuracy: 0.5505

Epoch 00020: LearningRateScheduler reducing learning rate to 0.009216.

Epoch 20/100

157/157 [=====] - 11s 71ms/step - loss: 0.5441 -
accuracy: 0.7916 - val_loss: 1.1338 - val_accuracy: 0.6275

Epoch 00021: LearningRateScheduler reducing learning rate to 0.009216.

Epoch 21/100

157/157 [=====] - 11s 71ms/step - loss: 0.5009 -
accuracy: 0.8090 - val_loss: 0.9677 - val_accuracy: 0.6515

Epoch 00022: LearningRateScheduler reducing learning rate to 0.009216.

Epoch 22/100

157/157 [=====] - 11s 71ms/step - loss: 0.4808 -
accuracy: 0.8197 - val_loss: 0.8817 - val_accuracy: 0.6960

Epoch 00023: LearningRateScheduler reducing learning rate to 0.009216.

Epoch 23/100

157/157 [=====] - 11s 71ms/step - loss: 0.4576 -
accuracy: 0.8256 - val_loss: 1.1543 - val_accuracy: 0.6415

Epoch 00024: LearningRateScheduler reducing learning rate to 0.008847359999999999.
Epoch 24/100
157/157 [=====] - 11s 71ms/step - loss: 0.3933 - accuracy: 0.8484 - val_loss: 1.5344 - val_accuracy: 0.5270

Epoch 00025: LearningRateScheduler reducing learning rate to 0.008847359999999999.
Epoch 25/100
157/157 [=====] - 11s 71ms/step - loss: 0.3917 - accuracy: 0.8475 - val_loss: 1.2672 - val_accuracy: 0.6385

Epoch 00026: LearningRateScheduler reducing learning rate to 0.008847359999999999.
Epoch 26/100
157/157 [=====] - 11s 71ms/step - loss: 0.3682 - accuracy: 0.8591 - val_loss: 1.0776 - val_accuracy: 0.6380

Epoch 00027: LearningRateScheduler reducing learning rate to 0.008847359999999999.
Epoch 27/100
157/157 [=====] - 11s 71ms/step - loss: 0.3532 - accuracy: 0.8661 - val_loss: 1.2599 - val_accuracy: 0.6070

Epoch 00028: LearningRateScheduler reducing learning rate to 0.008847359999999999.
Epoch 28/100
157/157 [=====] - 11s 71ms/step - loss: 0.3343 - accuracy: 0.8727 - val_loss: 0.9829 - val_accuracy: 0.6840

Epoch 00029: LearningRateScheduler reducing learning rate to 0.008847359999999999.
Epoch 29/100
157/157 [=====] - 11s 71ms/step - loss: 0.3237 - accuracy: 0.8749 - val_loss: 1.3291 - val_accuracy: 0.6375

Epoch 00030: LearningRateScheduler reducing learning rate to 0.008847359999999999.
Epoch 30/100
157/157 [=====] - 11s 71ms/step - loss: 0.3169 - accuracy: 0.8810 - val_loss: 1.1018 - val_accuracy: 0.6460

Epoch 00031: LearningRateScheduler reducing learning rate to 0.008847359999999999.
Epoch 31/100
157/157 [=====] - 11s 71ms/step - loss: 0.2684 - accuracy: 0.8983 - val_loss: 1.8160 - val_accuracy: 0.5205

Epoch 00032: LearningRateScheduler reducing learning rate to 0.008493465599999998.

Epoch 32/100


157/157 [=====] - 11s 71ms/step - loss: 0.2646 - accuracy: 0.8974 - val_loss: 0.9670 - val_accuracy: 0.6890


Training Time: 366.57504773139954 seconds

```
[19]: score5 = model5.evaluate(test_images,test_labels,verbose=0)
print('Test Loss - ',score5[0])
print('Test Accuracy - ',score5[1])
graphs(history5,score5)
```

Test Loss - 0.9669646620750427

Test Accuracy - 0.6890000104904175





Accuracy of the model: 68.90%

0.15 Results Discussion

Our Imagenet dataset consists of 8 classes with 1500 colored images in each class. For this dataset with the input size chosen as 128 x 128 with 3 channels the following results are obtained when trained until convergence.

Observations

- **VGG-16** - It is a complex deep CNN architecture with 16 weight layers and 3 fully connected layers. The model has a total of 64.5 million trainable parameters. The model produces an accuracy of 60.50 % over the test data.
- **VGG-19** - It is a variant of model 1 with 19 weight layers and 3 fully connected layers. It has a total of 70 million parameters. The model produces an accuracy of 59.64 % over the test data.
- **AlexNet** - This model is comparatively shallow than the other 2 models with only 5 weight layers and 3 fully connected layers. The models consists of 24.8 million trainable parameters and produces an accuracy of 68.50 % over the test data.
- **GoogLeNet** - Used inception module, it provided an accuracy of 69.74 % on the test data. This is the highest accuracy obtained on the test data.
- **ResNet** - Used residual module, it provided an accuracy of 58.85 % on the test data.
- **Combined Model 1** - It is a combination of AlexNet and Inception module. The model is shallow with two inception module. The test accuracy was 67.55 % the model was slightly overfitting, which was observed from the loss.
- **Combined Model 2** - It is a variation of combined model 1. Few regularizing parameters were added along with one more inception module, thus making it 3

inception modules. The test accuracy was 69.55 %, which was the second best accuracy observed on the test data.

Key Inference

- We can see that the more shallow AlexNet model has performed very well over the dataset, there are many reasons that attribute to this. It can be due to the fact that the images are considered at a lowered pixel ratio, in which some of the valuable features may be lost.
- We observe GoogleNet performed very well on the data. This is mainly attributed to the inception modules.
- The Combined model using the inception model, seemingly provided better results which can clearly supports the above stated claim.
- The relatively less performance of the VGG-19 model when compared with VGG-16 can be attributed to the depth of the CNN architecture. The input parameters are very less when considered with the actual depth of the model, as increased number of layers with less features to catch can reduce the activation of neurons with certain features and due to the implicit regularizations because of the depth of the model.
- Since the data set has only 1500 images for consideration and some of these images were seen to having obstructions such as human interventions, the model may be restricted to learn only certain features of the data, if trained on more data which is clean, the performance of the model can be considerably increased.
- Another potential factor that could be affecting the for model's performance is cropping the image. Although we crop the image from the lowest side in order to capture the maximum features out of it, in some cases the part of the image that is discarded may contain much more feature information than expected, which can be addressed by increasing the dataset size. By increasing the size of the data, these missed features may not affect the models performance as the probability that all the cropped images miss the same features is very less. #### Room for Improvements
- The current model only uses 1500 images per class to train. Increasing the datasets size would increase the accuracy to a very good extent.
- Increasing the input image size. Current Model uses 128 x 128 images, increasing the size would give more additional information to extract, thus increasing the performance.
- Using data augmentation, can also improve accuracy of the model.